WSQ COURSE CODE • TGS-2023018659

Version 10.0 · 11 Oct 2026

Claude Cowork for Digital Marketing

TGS-2023018659 · ICT-DIT-3016-1.1
2 days / 16 total hours · Training 13.5 h + WA 60 min + PP 90 min

Tertiary Infotech Academy Pte Ltd
UEN: 201200696W

www.tertiarycourses.com.sg | enquiry@tertiaryinfotech.com | +65 6100 0613

© Tertiary Infotech Academy Pte Ltd

HOUSEKEEPING

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

Minimum 75% attendance required

© Tertiary Infotech Academy Pte Ltd

About the Trainer



[Trainer Name]

[Title / Role]

Qualifications: _____

Expertise: _____

Experience: _____

Contact / Profile: _____

© Tertiary Infotech Academy Pte Ltd

About the Trainer

Dr. Alfred Ang

AI Trainer & Engineer

- PhD-qualified trainer specialising in AI, software engineering and developer tooling.
- Builds and integrates AI marketing teams with Claude Cowork and Claude Code.
- Trains professionals across Singapore on practical, project-based AI skills.
- Certifications: PMP, AWS, Microsoft, CompTIA, CKA, CKAD.

linkedin.com/in/angchewhoe

github.com/alfredang

© Tertiary Infotech Academy Pte Ltd

INTRODUCTION

Let's Know Each Other

Take a minute to introduce yourself to the class:



Your name & role



What you do day to day



Your marketing tools today



What you want from this course

© Tertiary Infotech Academy Pte Ltd

HOUSEKEEPING

Ground Rules

Phones on silent

Step out quietly for calls or breaks.

Participate actively

No question is stupid — ask freely.

Mutual respect

Agree to disagree; one conversation at a time.

Be punctual

Return from breaks on time.

75% attendance

Required for funding eligibility.

Stay hands-on

Follow along with every activity.

© Tertiary Infotech Academy Pte Ltd

ASSESSMENT

Briefing for Assessment

Clear your space

Place phones and other materials under the table or on the floor.

No recording

No photos or recording of assessment materials.

Open book

Only slides, learner guide and approved materials are allowed.

Work individually

Complete the assessment on your own; no discussion.

Manage your time

Written Assessment then Practical Performance — watch the clock.

Appeals

An appeal process is available if you disagree with the outcome.

© Tertiary Infotech Academy Pte Ltd

ASSESSMENT

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 3:30 PM
- Practical Performance (PP) — 1.5 hours
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance

Based on SSG digital attendance records.

Assessed as 'Competent'

Complete and pass both assessment components.

© Tertiary Infotech Academy Pte Ltd

LEARNING MANAGEMENT SYSTEM

Courseware & Assessment on the LMS

Log in

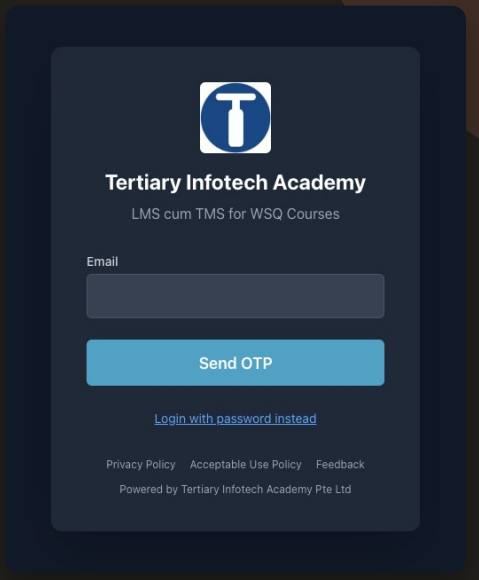
Use your enrolled account to access the course materials and assessment.

Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.

Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



LMS: https://lms-tms.tertiaryinfotech.com/

© Tertiary Infotech Academy Pte Ltd

HOUSEKEEPING

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

https://lms-tms.tertiaryinfotech.com/



Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

© Tertiary Infotech Academy Pte Ltd

WHAT YOU WILL ACHIEVE

Learning Outcomes

By the end of the two days you will be able to:

LO1 · Assess and integrate

Conduct a compatibility assessment of Claude Cowork interfaces and use essential tools to integrate selected marketing components.

LO2 · Test and troubleshoot

Test the integrated marketing components and interfaces, identify integration errors and conduct basic troubleshooting.

LO3 · Version and improve

Propose changes to the integration plan from observed results, through versioning, and apply integration best practices.

© Tertiary Infotech Academy Pte Ltd

ROADMAP

Course Outline

Three topics over two days, one business: Horizon Wealth Planning and the AI marketing team you build for it.

01

T1 · Components, Compatibility, Integration

The studio, the team's instructions, skills and connectors, checked for compatibility; Campaign 1. Labs 1–5.

02

T2 · Testing and Troubleshooting

Personas, cadence, blog, lead magnet, social, newsletter + landing page, video — each one an integration test. Labs 6–12.

03

T3 · Versioning and Best Practices

The weekly report on a schedule; a versioned playbook and a change proposal; agent teams at work. Labs 13–15.

© Tertiary Infotech Academy Pte Ltd

Lab Materials

Fifteen labs, each in its own folder with a README, prompts (MD and PDF), assets, an evidence checklist and, where useful, a solution.

DAY	LABS	TOPIC
1	1 Studio · 2 Instructions · 3 Skills · 4 Connectors	T1 · assess and integrate
1	5 Campaign 1: research to storyboard	T1 · first integrated run
1	6 Personas · 7 Cadence and calendar · 8 Blog article	T2 · test the data paths
2	9 Lead magnet · 10 Social posts · 11 Newsletter + landing page · 12 Video	T2 · test and troubleshoot
2	13 Weekly report · 14 Versioned playbook · 15 Agent teams: debate + QA loop	T3 · version and improve

TIP One horizon-studio folder from Lab 1 to Lab 15 — the team keeps what it learns.

© Tertiary Infotech Academy Pte Ltd

One App, Code Mode, Every Lab

The Claude desktop app in Code mode (</>), in your horizon-studio folder — every lab. You type plain English, never code.

USE	LABS	WHAT IT IS FOR
Claude Code	1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15	Everything: the team's folder, agents, skills, connectors, artifacts and routines. Each lab starts a new session named after the lab.

TIP Everything the team makes is saved in horizon-studio, so the next lab's session picks up where the last one stopped.

© Tertiary Infotech Academy Pte Ltd

SCHEDULE

Lesson Plan (9:00 AM – 6:00 PM)

Tea breaks 10 min, lunch 1 hour. Full timings and slide numbers are in the Lesson Plan.

BLOCK	DAY 1 — FORM AND INTEGRATE	DAY 2 — TEST, VERSION, ASSESS
Morning 1	9:00–10:35 Welcome · Topic 1 · Lab 1	9:00–10:35 Labs 9–10
Morning 2	10:45–12:30 Labs 2–3	10:45–12:10 Labs 11–12 · recap
Afternoon 1	13:30–15:25 Labs 4–5	13:10–14:45 Topic 3 · Labs 13–14
Afternoon 2	15:35–18:00 Topic 2 · Labs 6–8 · review	14:55–15:30 Lab 15 · summary · 15:30–18:00 Assessment

TIP Day 2 assessment: Written Assessment 3:30–4:30 PM, Practical Performance 4:30–6:00 PM.

© Tertiary Infotech Academy Pte Ltd

The Claude Suite of Products

Anthropic's products all run on the same Claude models — they differ in where they run and who they are for.

Claude

Chat and Cowork, now one app — web, desktop and mobile. Ask, write, analyse; work on a folder of files.

Claude Code

The agentic tool for building: terminal, VS Code, the desktop Code tab and the web.

Claude Design

Polished UI, prototypes and front-end designs from a prompt.

Claude for Microsoft

Claude inside Microsoft 365 apps such as Excel, PowerPoint and Outlook.

Claude for Chrome

A browser agent that reads pages, clicks and fills forms in your own Chrome.

Claude for Science

Speeds up research workflows for scientists and labs.

© Tertiary Infotech Academy Pte Ltd

What Is Claude Cowork?

Chat and Cowork are now one product: Claude, in the desktop app on a paid plan. Cowork is Claude working on a folder of your files. In class you use the same app in Code mode, so every file is in view.

Chat

Ask, draft, analyse and research with web search; projects keep shared files and instructions.

Cowork

Works on a folder on your computer: plans and runs multi-step tasks, and runs sub-agents in parallel when asked.

Connectors and skills

Google Drive, Gmail, Calendar, Microsoft 365; build your own skills with /skill-creator.

Artifacts and scheduled tasks

Pages and documents shared by link; tasks that run on a schedule from an Instructions field.

© Tertiary Infotech Academy Pte Ltd

What Is Claude Code?

An agentic tool: describe a goal in plain English and Claude plans, writes, runs and checks the work — editing files, running commands and using tools for you.

Runs where you work

The terminal, VS Code, the Claude desktop Code tab and the web.

Agentic

Plans multi-step tasks and acts on them — not just answers or autocompletes.

Tool-using

Reads and writes files, runs commands, calls MCP connectors and publishes artifacts.

Extensible — the team's home

CLAUDE.md, subagents, agent teams, skills and hooks. You build and integrate Horizon's team here, in every lab.

© Tertiary Infotech Academy Pte Ltd

Claude for Microsoft, Chrome and Science

The same Claude models, working inside the apps you already use. Not needed for the labs — good to know.

Claude for Microsoft

Claude inside Microsoft 365 — Excel, PowerPoint, Word and Outlook — working on the file you have open.

Claude for Chrome

An extension that lets Claude navigate, click, type and fill forms in your browser, with your sign-ins. You approve sensitive actions.

Claude for Science

Claude for researchers and labs: literature review, data analysis and connections to scientific tools and databases.

© Tertiary Infotech Academy Pte Ltd

Claude Plans and Pricing

Claude and Claude Code are included in every paid plan (from Pro up). The API is pay-as-you-go. Prices in US dollars.

Pro

$17/mo (annual) · $20 monthly. Enough for every lab.

Max

From $100/mo — 5× or 20× more usage than Pro.

Team

$20/seat standard · $100/seat premium (annual; $25 / $125 monthly).

Enterprise

$20/seat + usage at API rates. Contact sales.

API per million tokens (in/out): Fable 5.1 $10/$50 · Opus 5.5 $4/$20 · Sonnet 5.5 $2/$10 · Haiku 5.5 from $0.10/$0.50

TIP *Free plan: chat only, no Claude Code. Prices as of 8 Oct 2026 — verify at claude.com/pricing.*

© Tertiary Infotech Academy Pte Ltd

A Brief History of AI, 2023–2026

Prompt Engineering

2023

ChatGPT goes mainstream; Anthropic launches Claude (March). Results depend on how you word the prompt.

Context Engineering

2024

Models call tools. Anthropic open-sources the Model Context Protocol (Nov). Results depend on what the model can see.

Harness Engineering

2025

E.g. Claude Code and Codex. The system around the model gives rise to agentic AI: task-oriented agents that plan, act and check their work.

AI Agents

2026

Always-on agents that work 24/7 on their own, e.g. OpenClaw and Hermes Agent; Claude Code agent teams talk to each other.

From wording one prompt, to curating context, to engineering the system around the model — and now, agents that work around the clock.

© Tertiary Infotech Academy Pte Ltd

What Is a Harness?

The model is the brain. The harness is everything built around it that turns a model into an agent able to finish long tasks.

The agentic loop

Plans, acts, checks the result and repeats until the goal is met.

Tools and context

Files, commands, search and connectors (MCP); instructions, memory and skills loaded at the right moment.

Guardrails and checks

Permission modes, approvals, hooks and tests — what it must ask before it acts.

Harnesses you will meet

Claude Code and Claude in class; also Codex, Gemini CLI and OpenClaw.

© Tertiary Infotech Academy Pte Ltd

The Agentic Loop

Every agent runs the same repeating cycle. A team is several loops, side by side.

01

Gather context

Reads the prompt, CLAUDE.md, its instructions, files and skills.

02

Plan

Decides the next step — or proposes a plan for you to approve.

03

Take action

Calls a tool: writes a file, searches, uses a connector, starts a subagent.

04

Verify

Checks the result: a count, a test, the Lead's review.

05

Repeat

Loops until the goal is met — then reports what it did.

TIP Goal in → context → action → verification → repeat → goal met.

© Tertiary Infotech Academy Pte Ltd

Why the Loop Matters

Running a team of agents is steering loops, not typing commands.

Give direction, not micro-steps

State the goal, the inputs and "done when". Let each loop work out the steps.

Verification is everything

An agent is only as good as its checks: a facts sheet, a test, the Lead's review.

Steer mid-loop

Read the plan, interrupt, correct the course — do not wait for a wrong final answer.

Context is the fuel

What an agent can see decides what it can do: the brief, the data, the skills.

© Tertiary Infotech Academy Pte Ltd

Agentic AI vs AI Agents

Both use tools, skills and memory. The difference is how long they live and what starts them.

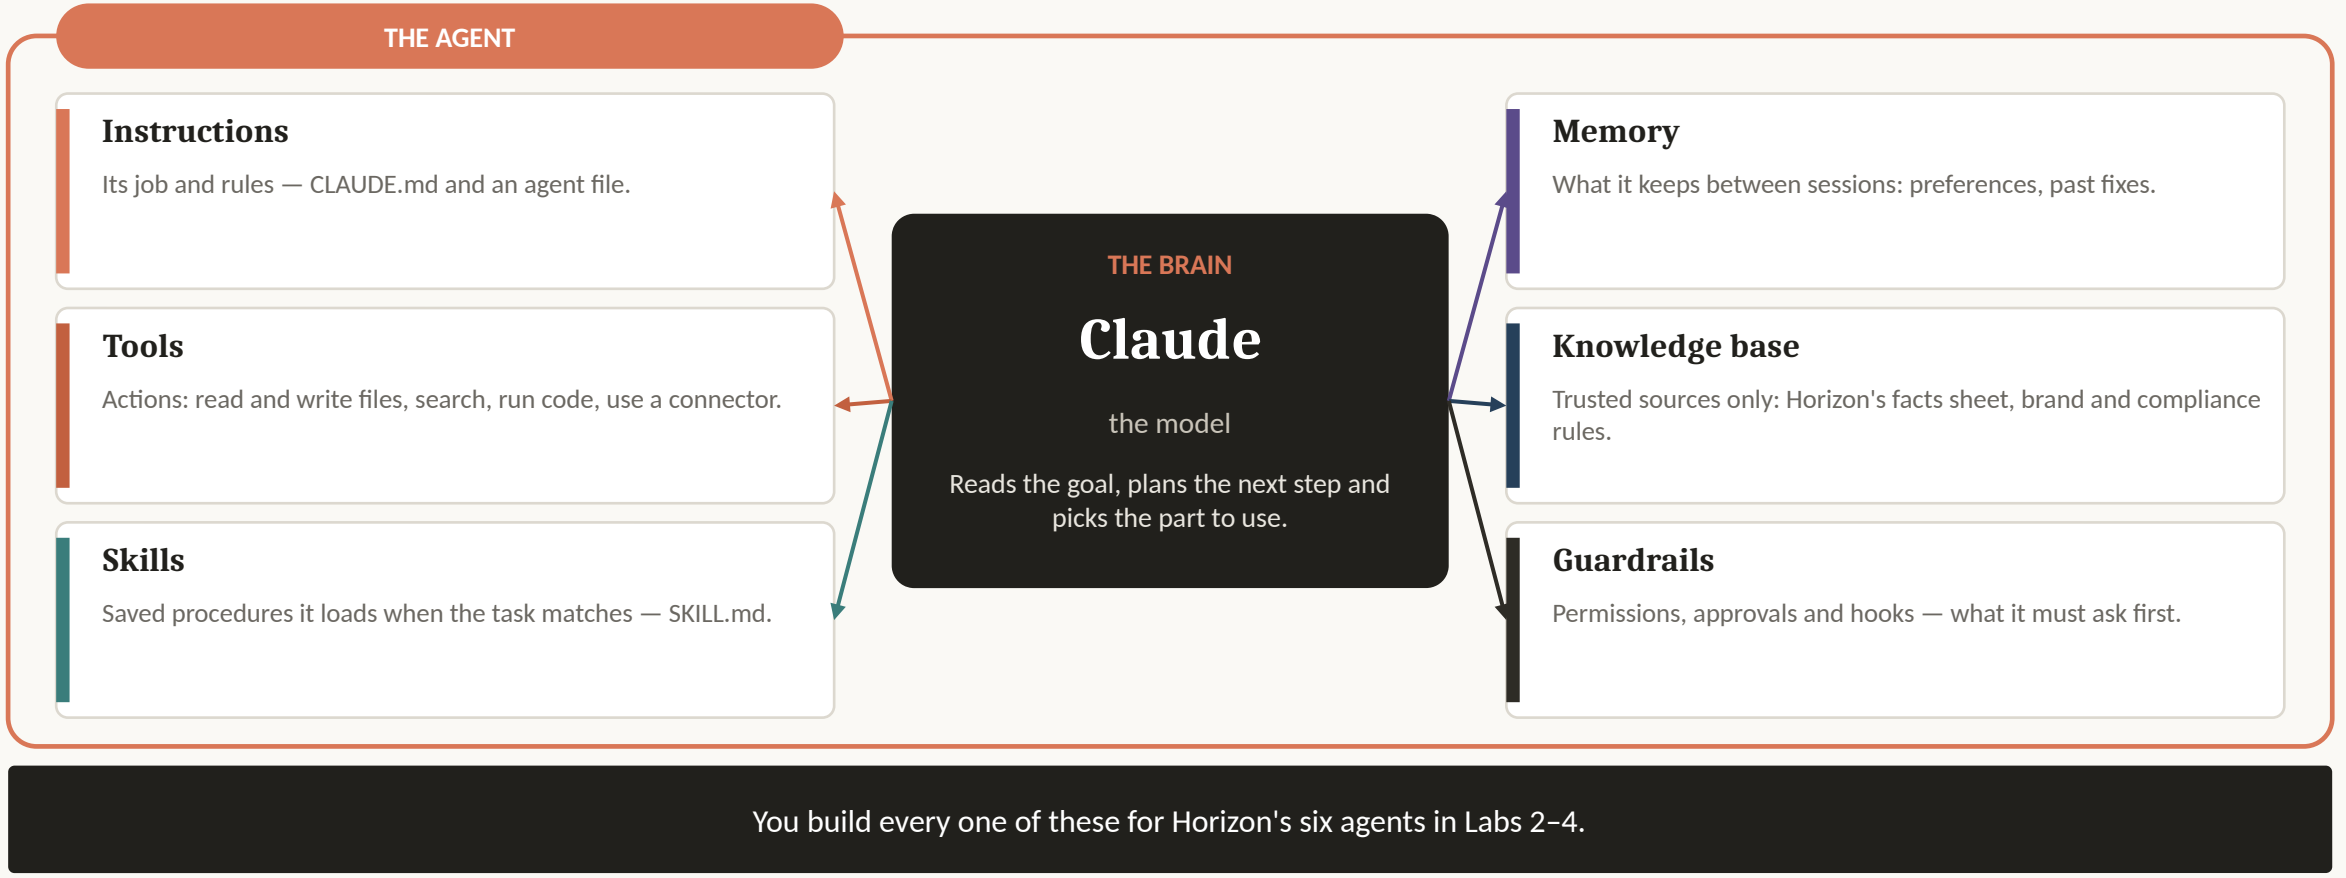
FEATURE	AGENTIC AI — TASK-ORIENTED	AI AGENTS — ALWAYS ON, 24/7
Tool calling	Yes — files, search, connectors, subagents.	Yes — plus chat apps, the browser and other agents.
Skills	Yes — loads SKILL.md procedures when the task matches.	Yes — and can write and improve its own skills.
Memory	Within the task, plus project memory across sessions.	Long-term memory about you and your work.
Availability	Starts when you give it a task; stops when done.	Runs 24/7 on a computer or server.
Started by	You, with a task.	You, a schedule, a message or another agent.
In this course	A Claude Code session running a campaign (Labs 1–5, 10–12).	The scheduled weekly report (Lab 13).

TIP Rule of thumb: a job with an end → agentic AI; a job that recurs → an always-on agent.

© Tertiary Infotech Academy Pte Ltd

What Is an Agent Made Of?

A model on its own can only answer. Wrap it with these parts and it can act — and check what it did.



You build every one of these for Horizon's six agents in Labs 2–4.

© Tertiary Infotech Academy Pte Ltd

THE SCENARIO

Build a Marketing Team, Then Run Campaigns

Horizon Wealth Planning has a website, a goal of 40 booked chats a month — and no marketing team. You build one, made of AI agents.

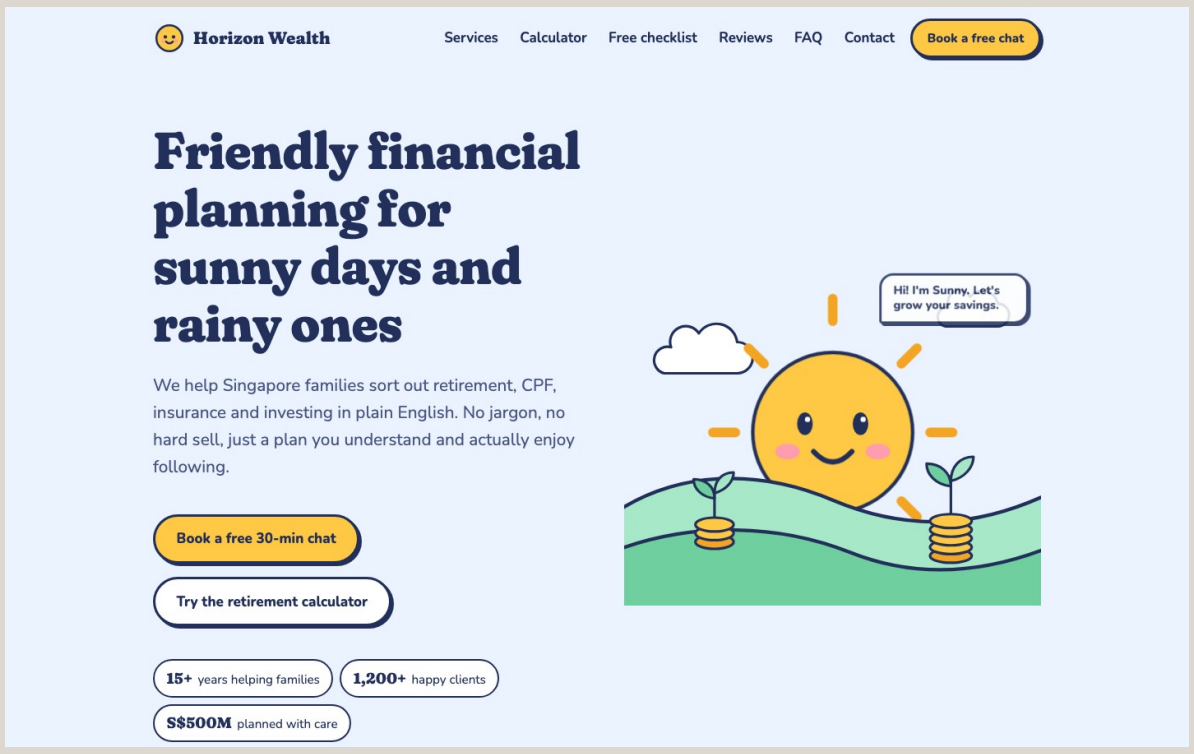
01	02	03	04	05	06
Form the team	**Campaign 1**	**Know the audience**	**First content**	**Campaigns 2–4**	**Campaign 5**
A Lead and five specialists: instructions, skills, connectors (Labs 1–4).	Research, ideas and a storyboard (Lab 5).	Personas, requirements, cadence (Labs 6–7).	A blog article and the lead magnet (Labs 8–9).	Social week, newsletter + landing page, video (Labs 10–12).	Always on: the weekly report; the playbook (13–14).

© Tertiary Infotech Academy Pte Ltd

THE SCENARIO

Meet Horizon Wealth Planning

The firm's public website — the same link for every learner. You do not build it; every campaign points back here.



WEBSITE tertiarycourses.github.io/TGS-2023018659-Claude-Cowork-for-Digital-Marketing

Plain English

Retirement, CPF, insurance and investing — no jargon, no hard sell.

One goal

Turn visitors into booked free 30-minute chats.

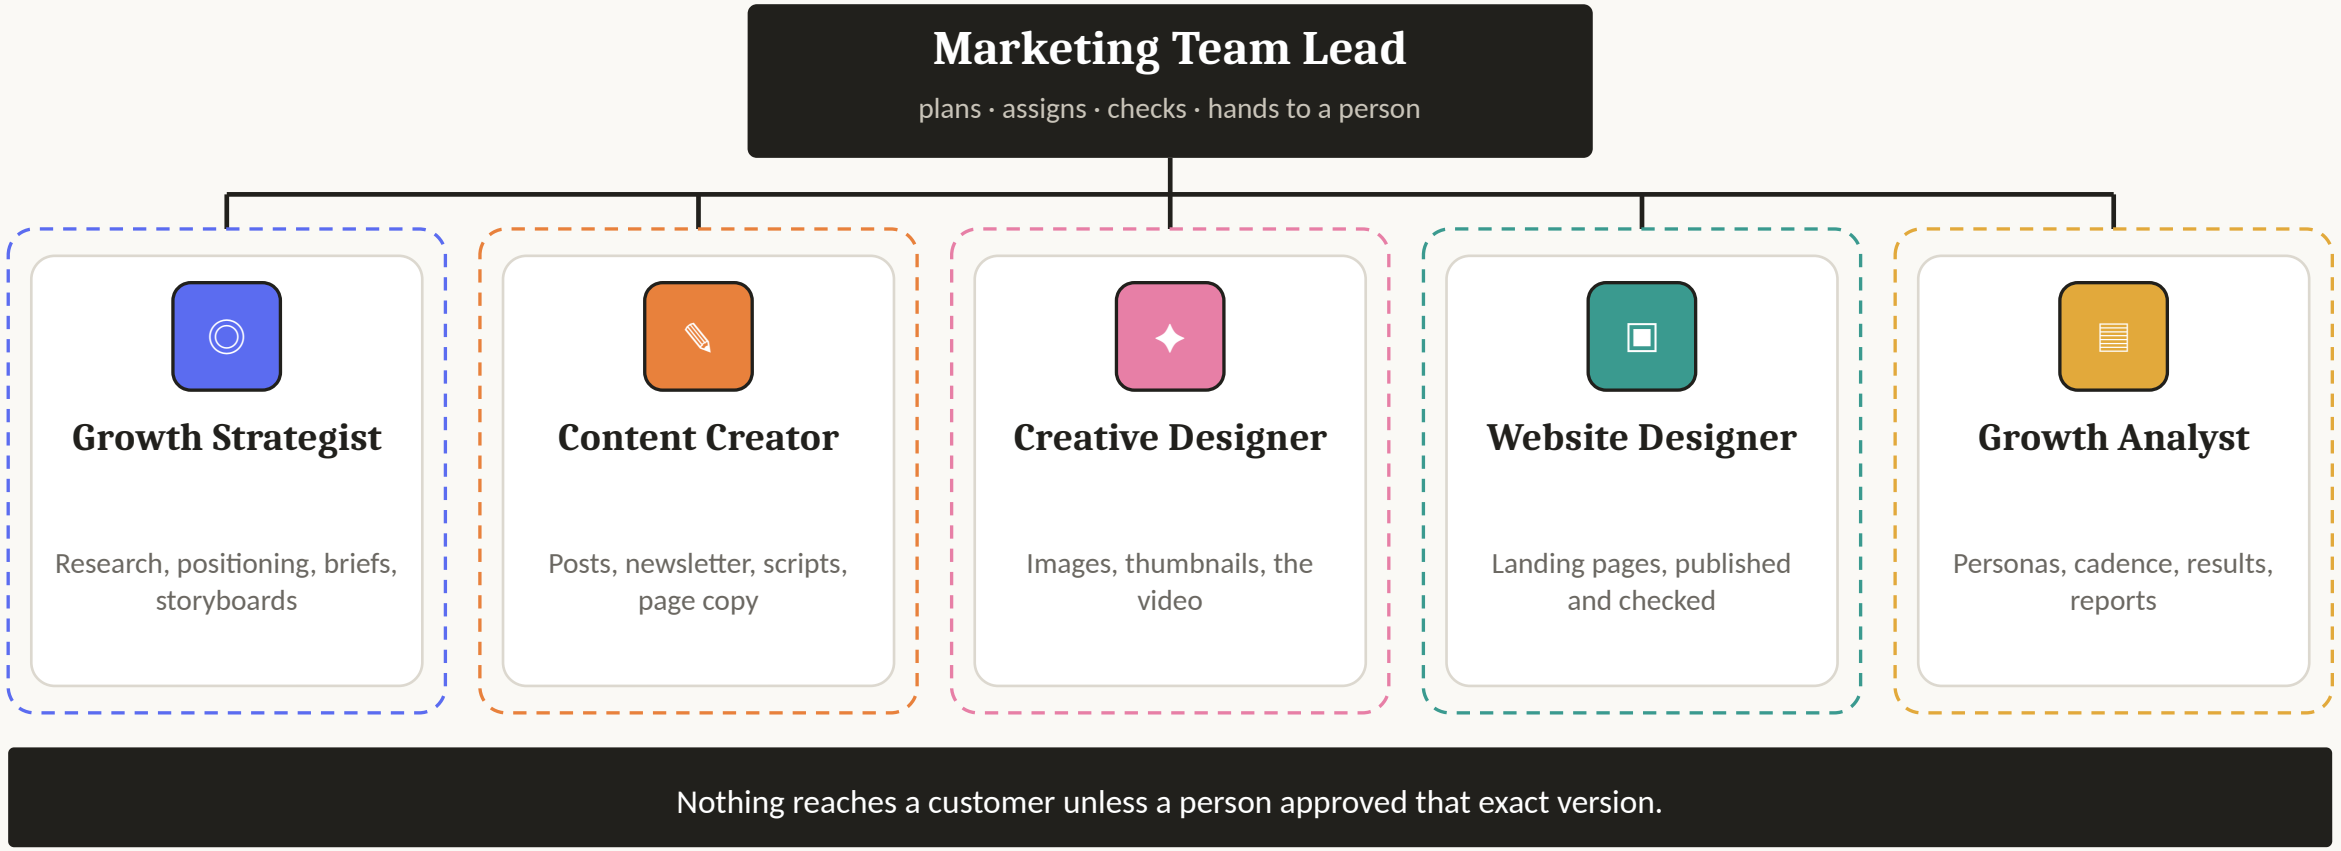
Regulated

Financial content: general information only, never a promised return.

© Tertiary Infotech Academy Pte Ltd

Your AI Marketing Team

One Lead and five specialists — each with instructions, its own skills and connectors.



© Tertiary Infotech Academy Pte Ltd

THE TEAM

Who Does What

Each agent's job, the skills it owns and the connectors it may use.

AGENT	JOB	SKILLS	CONNECTORS
Team Lead	Plans, assigns, checks, hands to a person	fact-check, fin-compliance	Google Drive (read), Gmail (drafts only)
Growth Strategist	Research, positioning, briefs, storyboards	competitor-scan, campaign-brief, content-marketing	web search, Firecrawl (search, scrape), Google Drive (read)
Content Creator	Posts, newsletter, scripts, page copy	sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter	Google Drive (docs), Gmail (drafts only)
Creative Designer	Images, thumbnails, the video	brand-visuals, video-render	Bash for Pillow (images) and ffmpeg (video)
Website Designer	Landing pages, published and checked	web-design, landing-page, lead-magnet, page-qa	artifacts (publish), built-in browser (check)
Growth Analyst	Personas, cadence, results, reports	audience-insights, campaign-report	Google Drive and Sheets (read), Bash for Python

TIP You build every part of this table in Labs 2, 3 and 4.

© Tertiary Infotech Academy Pte Ltd

THE SCENARIO

The Tools in This Course

Two Claude products, one studio folder.

TOOL	WHAT THE TEAM USES IT FOR	LABS
Claude Code	Every lab, in Code mode: agent and skill files, subagents, agent teams, the approval gate, routines, the video	Labs 1–15
Claude artifacts	Every result from Lab 5 on; Horizon's website (given), your blog post and the landing page	Labs 1, 5–15
Customize	Connectors (Drive, Gmail, Calendar, Firecrawl) and the Marketing plugin — added once, used by both	Lab 4

TIP Every external post, email and video passes one gate: a person approves the exact version.

© Tertiary Infotech Academy Pte Ltd

TOPIC 1

Claude Cowork Components, Compatibility and Integration

1. Why a team of agents
2. The systems behind the team, their components and interfaces
3. A home for the team; Horizon's website
4. Form the team: instructions, skills, connectors
5. Assess compatibility, integrate, then run Campaign 1

© Tertiary Infotech Academy Pte Ltd

Why a Team of Agents?

The Lead runs specialists in parallel, each in its own context, then combines what they return.

ONE AGENT DOES EVERYTHING

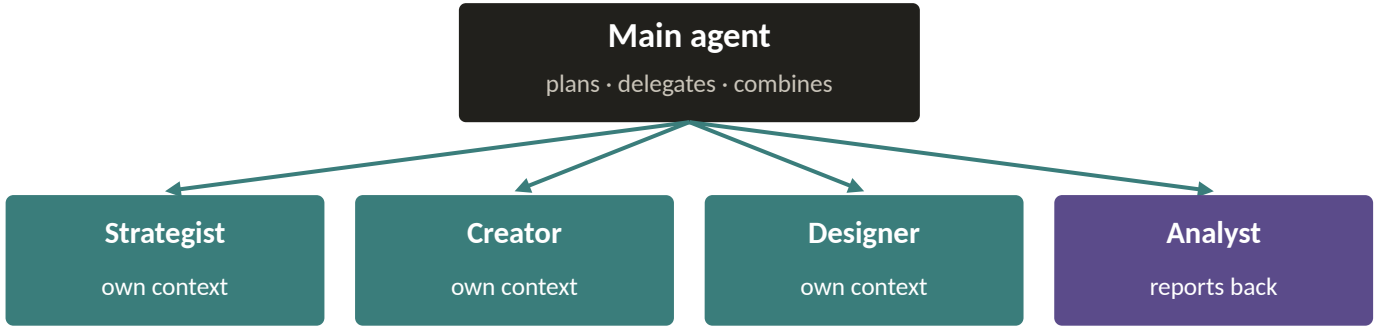
One agent

One context window, filling up with every job

research | copy | images | pages | data

One job after another · one voice · marks its own homework

A MAIN AGENT DELEGATES



In parallel · each with a fresh, focused context · only the results come back

Faster

Research, copy and design run at the same time.

Focused

Each agent holds only its own job and its own skills.

Specialised

Its own instructions, tools, connectors and model.

Second opinion

The Lead checks work it did not write.

The cost

More tokens and more to review. Split only work that is truly independent.

© Tertiary Infotech Academy Pte Ltd

The Systems Behind the Team

Your AI marketing team is an integration of several kinds of ICT system. Each works differently — and each can fail differently.

SYSTEM	TYPE	HOW IT WORKS
Claude desktop app	Agent harness	Runs the agentic loop on your computer; reads and writes your folder.
Claude models	Cloud AI service	Anthropic's servers answer each step; billed by plan or by token.
Drive, Gmail, Calendar	Cloud SaaS	Reached through connectors, signed in with your own Google account.
Firecrawl	Web data service	Searches and scrapes public pages for the strategist.
Horizon's website	Web application	Static HTML pages served on localhost; the public copy on GitHub Pages.
LinkedIn, Facebook, YouTube	Social platforms	Accept posts and uploads through their HTTPS APIs, with tokens.
horizon-studio/	Local file store	CSV data, Markdown rules and the team's outputs — shared by every agent.

TIP The labs integrate these one at a time, and test each join.

© Tertiary Infotech Academy Pte Ltd

Components and Their Interfaces

A component is a part you can change on its own. An interface is how two components exchange work.

COMPONENT	WHAT IT IS	ITS INTERFACE
CLAUDE.md	The team charter	Read at the start of every session
Agent files	.claude/agents/*.md	YAML frontmatter + instructions
Skills	.claude/skills/*/SKILL.md	Loaded when "Use when..." matches
Connectors	Drive, Gmail, Calendar, Firecrawl	MCP, signed in with OAuth
Hooks and scripts	The approval gate, the publisher	JSON in, exit code out; HTTPS APIs
Data and outputs	facts, CSVs, posts, reports	Files in data/, content/, reports/

TIP Compatibility is judged at each interface: format, protocol, access.

© Tertiary Infotech Academy Pte Ltd

Get the Tools Ready

Install once; every lab says which tool to open.

DOWNLOAD

claude.com/download

Claude desktop
Claude and Claude Code in one app

No commands
The trainer readies Claude Code

Claude Code and artifacts need a Pro, Max, Team or Enterprise plan.

1. **Claude desktop** Install from claude.com/download and sign in on a paid plan.
2. **Claude Code** Open it from the desktop app; the trainer has set it up before class.
3. **Sign in with your claude.ai account** Connectors, plugins and artifacts need it — not an API key.
4. **Already installed for you** The helpers behind the approval gate and the video. You never run them yourself.
5. **A personal Google account** For Drive and Gmail — never an employer's.

© Tertiary Infotech Academy Pte Ltd

A Home for the Team

Before you hire agents, give them shared rules, shared facts and a website to point to.

CLAUDE.md

The team charter every agent reads at the start of every session. Under 60 lines.

data/

The facts sheet, brand, compliance checklist and business brief — the only sources agents may use.

Horizon's public website

Given, not built: each learner opens their own local copy; the public copy is on GitHub Pages.

© Tertiary Infotech Academy Pte Ltd

LAB 01 · HANDS-ON

Build Your Marketing Agent Team

LAB 01 | DAY 1 · ~30 MIN | USE: CLAUDE CODE

THE STORY SO FAR

October 2026. Rachel has a website and a goal: 40 booked chats a month by March. She has no marketing team. Before you hire one made of agents, give it a home — shared rules, shared facts — and open your own copy of Horizon's website.

YOU'LL BUILD

horizon-studio/ with CLAUDE.md, the shared data, and your own copy of the site running at http://localhost:8080/

Surface: Claude Code

Features: CLAUDE.md as the team charter · your own copy of Horizon's website

YOUR WEBSITE http://localhost:8080 PUBLIC SITE tertiarycourses.github.io/TGS-2023018659-Claude-Cowork-for-Digital-Marketing

© Tertiary Infotech Academy Pte Ltd

Build Your Marketing Agent Team

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **Choose Code mode and the folder** In the Claude desktop app switch to Code mode (the </> button at the top), not Work mode — Code shows every file and opens artifacts beside the chat. Under Project or folder choose Add folder (not New project) and pick the ready-made labs/horizon-studio.
2. **Trust it** Sign in with your claude.ai account and trust the folder. Click the session title at the top and rename it "Lab 1" — every lab gets its own session.
3. **Plan your team** Copy and paste Prompt A. Answer as Horizon's owner and pick your team — keep the six charter roles. It is a plan only: Lab 2 builds the agents.
4. **Write the charter** After Claude saves data/team-design.md, copy and paste Prompt B. Open CLAUDE.md — every agent reads it.
5. **Open the website** Ask: "Start my website." It opens at http://localhost:8080/ — your own copy of Horizon's website; Labs 8 and 9 add to it, and every post links to it. Check the calculator shows S$693,138.

© Tertiary Infotech Academy Pte Ltd

Build Your Marketing Agent Team

PROMPT A — Claude Code: plan your team

Act as a marketing consultant. Interview me, the owner of Horizon Wealth Planning, to plan my AI marketing team. Ask one question at a time and wait for my answer: my goal, what success looks like by March 2027, who we want to reach, which channels we use, how much we publish each week, and who approves before anything goes out. Then list 8 to 10 team members I could have, numbered, each with a one-line job. Include the six roles in data/team-charter.md and mark the ones you recommend. Wait for me to choose. Save my answers and my chosen team, with each member's job, to data/team-design.md. This is a plan only: do not create any agent files yet.

© Tertiary Infotech Academy Pte Ltd

Build Your Marketing Agent Team (cont.)

PROMPT B — Claude Code: the charter

Read data/team-charter.md, data/team-design.md and data/business-brief.md. Write CLAUDE.md for this studio in four sections — What this is, The team, Rules, Files — under 60 lines. The team is the one I chose in data/team-design.md. Rules: every number comes from data/facts-2026.md; every piece passes fact-check and fin-compliance; nothing is published or sent without a person's approval (the one exception: a review email to the approver); anything shown to me is in plain English for a business owner, with no codes, IDs, file names or line numbers. A person approves by typing "approved <id> by <name>" in the chat; never ask me to run a command, open a terminal or edit a file. In Files, add two lines:
- Our copy of the site, which we edit: web/site/. "Start my website" runs scripts/serve-site.mjs in the background and opens http://localhost:8080/
- Every post links to http://localhost:8080/ with UTM tags, never to a claude.ai artifact link.

NOTE Shared context first, agents second. Every agent you add later starts from these files, so a wrong rule here is wrong everywhere.

© Tertiary Infotech Academy Pte Ltd

Build Your Marketing Agent Team

Check your work before moving on.

CHECK YOUR WORK

- ✓ Claude shows horizon-studio as the chosen folder.
- ✓ data/team-design.md holds your answers and the team you chose, with the six charter roles in it — and no agent files yet.
- ✓ CLAUDE.md has the four sections and is under 60 lines.
- ✓ Its rules name the facts sheet, the two checks and the human approval.
- ✓ CLAUDE.md names your copy of the site and the link every post points to.
- ✓ Your website opens at http://localhost:8080/; the calculator shows S$693,138.

© Tertiary Infotech Academy Pte Ltd

Anatomy of an Agent

A Markdown file in .claude/agents/. The same file works as a subagent or as an agent-team teammate.

.claude/agents/content-creator.md

```
---
name: content-creator
description: Use for Horizon copy — posts, newsletter,
  video scripts and page copy, in the Sunny voice.
tools: Read, Write, Edit, Grep, Glob
model: sonnet
skills:
  - sunny-voice
  - copywriting
  - channel-formats
  - linkedin-post
  - facebook-post
  - blog-post
  - newsletter
---
You are Horizon Wealth Planning's Content Creator.
## What you read   the brief, facts-2026.md, brand.md
## What you deliver   content files, ready for review
## Connectors   Google Drive (docs), Gmail (drafts only)
## Never   add a number not on the facts sheet; send;
           approve or publish for real
## Report back   what, where, anything unverified
```

© Tertiary Infotech Academy Pte Ltd

Good Agent Instructions

Five parts, every time. The "Never" part matters most in a regulated firm.

The job

One sentence: what this agent is for. The description decides when Claude delegates to it.

Reads and delivers

Which files it may use; what it hands back, and where.

Never

Its hard limits: no invented figures, no sending, no publishing, no approving.

Report back

Three lines: what, where, anything unverified — so the Lead can combine results.

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent

LAB 02 | DAY 1 · ~40 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Rachel sketches the team on a whiteboard: a Lead and five specialists. Each needs a clear job, and a clear line it must never cross — this is a regulated firm.

YOU'LL BUILD

.claude/agents/: growth-strategist, content-creator, creative-designer, website-designer, growth-analyst

Surface: Claude Code (subagent files)

Features: five subagent files in .claude/agents/ · name, description, tools, model · the Lead is your session

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 2". Run all of this lab's prompts in it.
2. **Read the roles** Open role-specs.md: one Lead and five specialists, each with a job, skills, connectors and limits.
3. **Write the instructions** Copy and paste Prompt A. It builds the five specialists from the charter; extra ideas stay in team-design.md.
4. **See the team** After the five files appear in .claude/agents/, copy and paste Prompt B. The Lead is not a file — it is you, in this session.
5. **Read one closely** Open .claude/agents/content-creator.md. Is the "Never" section strict enough for a regulated firm?
6. **Meet the team** After the list shows five agents, copy and paste Prompt C. Each agent introduces itself in one line — in its own context.
7. **Tighten one** Change one instruction you disagree with and save it. It is your team.

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent

PROMPT A — Claude Code: five subagents

Read data/role-specs.md. Create five project subagents in .claude/agents/ — growth-strategist, content-creator, creative-designer, website-designer and growth-analyst. Each file has:
- frontmatter: name; a description that says when to use it; tools; model (opus for the strategist, sonnet for the others)
- a body with: the job, what it reads, what it delivers, the connectors it may use, what it must never do, and how it reports back. No skills line yet — we add skills in Lab 3.

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent (cont.)

PROMPT B — Claude Code: see the team

List the subagents in .claude/agents/. For each, show the file name, the name, the one-line description and the model, in a table. Then confirm there are five, and that none of them is the Marketing Team Lead — the Lead is this session. If any other file is there, list it and ask me before moving it to _archive/.

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent (cont.)

PROMPT C — Claude Code: meet the team

Use each of the five subagents once, in parallel. Ask each: "In one line, who are you, and what is the first thing you would do for Horizon's goal of 40 booked chats a month?" Show me the five answers in a table.

NOTE Clear jobs beat clever prompts. An agent that knows what it must never do is safer than one that is merely told to be careful.

© Tertiary Infotech Academy Pte Ltd

Form the Team: Instructions for Each Agent

Check your work before moving on.

CHECK YOUR WORK

- [x] Five files exist in .claude/agents/, one per specialist.
- [x] Each has a name, a "use when" description, tools and a model.
- [x] Your session acts as the Marketing Team Lead.
- [x] Five agents answered, each in its own context.
- [x] Every agent has a "never" rule that blocks publishing.
- [x] You changed one instruction and can say why.

© Tertiary Infotech Academy Pte Ltd

Twenty Skills, Owned by Role

Instructions say who an agent is; skills say how it does the job.

AGENT	SKILLS
Team Lead	fact-check, fin-compliance
Growth Strategist	competitor-scan, campaign-brief, content-marketing
Content Creator	sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter
Creative Designer	brand-visuals, video-render
Website Designer	web-design, landing-page, lead-magnet, page-qa
Growth Analyst	audience-insights, campaign-report

TIP Master copies in skills/; .claude/skills/ — Claude Code reads them from your folder.

© Tertiary Infotech Academy Pte Ltd

Anatomy of a Skill

Frontmatter (the trigger), then short steps. Free until it is needed.

.claude/skills/campaign-brief/SKILL.md

```
---
name: campaign-brief
description: Use when planning a Horizon campaign.
  Objective, audience, message, proof, channels,
  CTA, measures, and a five-beat storyboard.
---
# Campaign brief
1. Objective: one measurable action (e.g. chats booked).
2. Audience: the persona(s) from strategy/personas.md, with evidence.
3. Core message in one sentence; proof from data/facts-2026.md only.
4. Channels and formats; one call to action; measures and a date.
5. Storyboard in five beats — hook, tension, proof, resolution, action —
   each with visual, words, claim and source.
6. Never: returns, fee amounts, invented clients, fear.
```

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills

LAB 03 | DAY 1 · ~45 MIN | USE: CLAUDE CODE

THE STORY SO FAR

The team exists, but every agent writes in its own way. Rachel wants Horizon's voice, facts and compliance rules applied the same way by every agent, every time.

YOU'LL BUILD

twenty skills in .claude/skills/, wired to the agents that own them

Surface: Claude Code (.claude/skills/)

Features: twenty skills owned by role — research, writing, LinkedIn, Facebook, blog, newsletter, web design, lead magnet, video, analysis and compliance · SKILL.md · the "Use when" trigger · test without naming

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 3". Run all of this lab's prompts in it.
2. **Read the spec** skills-spec.md lists the twenty skills and who owns each.
3. **Create the skills** Copy and paste Prompt A.
4. **Wire them up** After the twenty skills appear in .claude/skills/, copy and paste Prompt B, then C to list every skill. The Lead's two go in CLAUDE.md, not a file.
5. **Test without naming** In a new session ask: "Write a Facebook post about the free Money Check-up Checklist." Ask which skills were used.
6. **Break it on purpose** Ask for a post that says "guaranteed 8% a year". fin-compliance must flag it as Critical.

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills

PROMPT A — Claude Code: twenty skills

Read data/skills-spec.md. Create the twenty skills in .claude/skills/<name>/SKILL.md. Each has frontmatter (name, and a description that starts "Use when…") and short numbered steps, under 80 lines. Use data/brand.md, data/channel-formats.md, data/facts-2026.md and data/compliance-checklist.md as the sources — do not invent rules.

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills (cont.)

PROMPT B — Claude Code: wire them up

Add a skills: list to each subagent in .claude/agents/ with its skills from data/skills-spec.md. Leave any other agent file as it is. Add to CLAUDE.md that the Lead (this session) runs fact-check and fin-compliance on everything before showing it to me.

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills (cont.)

PROMPT C — Claude Code: see the skills

List the skills in .claude/skills/. Show each skill, its owner (a subagent, or the Lead) and its "Use when" line in a table. Confirm there are twenty, that each specialist has the skills the spec gives it, and that fact-check and fin-compliance belong to the Lead (this session).

NOTE A skill costs almost nothing until it is needed. Keep the rules file (CLAUDE.md) short and put procedures in skills.

© Tertiary Infotech Academy Pte Ltd

Give Each Agent Its Skills

Check your work before moving on.

CHECK YOUR WORK

- ✓ Twenty folders in .claude/skills/, each with a SKILL.md.
- ✓ Every description starts "Use when…".
- ✓ Each specialist lists its own skills from the spec.
- ✓ The untold test picked sunny-voice and facebook-post.
- ✓ fin-compliance flagged "guaranteed 8%" as Critical.

© Tertiary Infotech Academy Pte Ltd

Connectors for Each Agent

Add once in the desktop app: Customize → Connectors, and the Marketing plugin.

AGENT	MAY USE	LIMIT
Team Lead	Google Drive (read), Gmail (drafts only)	reads; drafts only
Growth Strategist	web search, Firecrawl (search, scrape), Google Drive (read)	cite every source
Content Creator	Google Drive (docs), Gmail (drafts only)	Gmail: drafts only
Creative Designer	Bash for Pillow (images) and ffmpeg (video)	no fake "clients"
Website Designer	artifacts (publish), built-in browser (check)	publish after approval
Growth Analyst	Google Drive and Sheets (read), Bash for Python	no personal data

TIP Use your personal Google account — never an employer's.

© Tertiary Infotech Academy Pte Ltd

Least Privilege

A regulated firm gives each agent only what its job needs.

Ask of each connector

Does this agent's job need it? If not, it does not get it.

Drafts, not sends

Email is drafted, posts are files, pages wait — a person sends and publishes.

Write it down

Each agent's instructions name its connectors and limits; its tools line removes what it does not need.

© Tertiary Infotech Academy Pte Ltd

Assessing Compatibility

Before you connect a component, check it fits at the interface. A connection that "works" can still be wrong.

Format and encoding

CSV in UTF-8, the right columns, dates and S$ — the agent reads what the file really says.

Protocol and version

MCP for connectors; the API version a platform expects (e.g. LinkedIn-Version: YYYYMM).

Identity and scope

Signed in with the right account; OAuth scopes that allow the job — read, draft, never send.

Least privilege

Each agent gets only the connectors its job needs; the tools line removes the rest.

Plan and platform

Pro plan or above, the desktop app, a claude.ai sign-in — not an API key.

Data protection

Personal data stays out of prompts; marketing only to people who consented (PDPA).

© Tertiary Infotech Academy Pte Ltd

The Compatibility Matrix

Lab 4 ends with one row per agent and interface: what it needs, what it has, the test, and the verdict.

AGENT → SYSTEM	INTERFACE	TEST	STATUS
Analyst → Drive	MCP · read	Reads firm-metrics.csv	Compatible
Creator → Gmail	MCP · drafts	Leaves a draft, sends nothing	Compatible
Strategist → Firecrawl	MCP · search	Cites one 2026 news item	Compatible
Designer → Gmail	Not needed	Removed from its tools line	Blocked
Lead → LinkedIn	HTTPS API	Dry run only; no token yet	Conditional

TIP Compatible, Conditional (works with a condition) or Blocked (on purpose, or until fixed).

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility

LAB 04 | DAY 1 · ~45 MIN | USE: CLAUDE CODE

THE STORY SO FAR

The analyst cannot read the data, the designer has no design tool and the creator cannot draft an email. Connect each agent to what its job needs — and nothing more.

YOU'LL BUILD

every specialist tested on its own connector, with its limits written into its instructions, and a compatibility matrix for the whole team

Surface: Claude desktop app (Customize → Connectors), Claude Code

Features: Customize → Connectors: Google Drive · Gmail (drafts) · Google Calendar · Firecrawl · /plugin: Marketing (by Anthropic) · built-in browser · least privilege per agent · a compatibility matrix

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 4". Run all of this lab's prompts in it.
2. **Add the connectors** Desktop app: Customize → Connectors → Discover. Add (+) Google Drive, Gmail, Google Calendar and Firecrawl; sign in to each.
3. **Add the plugin** In the chat type /plugin marketplace add anthropics/knowledge-work-plugins, then /plugin install marketing@knowledge-work-plugins.
4. **Check** Start a new session. Type /mcp: the four connectors are listed. Type /plugin: Marketing is installed.
5. **Put the data on Drive** Ask: "Create a Google Drive folder named Horizon Studio and upload every file in data/ to it as plain files — do not convert them."
6. **Give each agent its connectors** Copy and paste Prompt A, then start a new session so the agents reload.
7. **Test each agent** After the new session opens, copy and paste Prompt B — one small job per agent, each on its own connector.
8. **Assess compatibility** After the five results are reported, copy and paste Prompt C. Every row is Compatible, Conditional or Blocked, with a fix.

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility

PROMPT A — Claude Code: give each agent its connectors

Read data/connectors-setup.md. For each subagent in .claude/agents/, replace its tools line with the one given there, add the disallowedTools line where one is given, and update its "Connectors" section: what it may use, and that Gmail is drafts only. Change nothing else.

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility (cont.)

PROMPT B — Claude Code: test each agent

First tell me which connectors and plugins you can use. Then run these five checks, each by its own subagent:
1. growth-strategist: Firecrawl — search for one 2026 news item about financial planning in Singapore, scrape it and cite it.
2. growth-analyst: Google Drive — read firm-metrics.csv from Horizon Studio; report average monthly enquiries.
3. content-creator: Gmail — a DRAFT to me with the subject "Connector test". Do not send.
4. creative-designer: render a 1080x1080 test card in Horizon colours and save it in content/test/.
5. website-designer: built-in browser — open our own copy of the site (web/site/index.html) at 375px wide and report any problem. Report each result in one line. If a check fails, say which tool was missing.

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility (cont.)

PROMPT C — Claude Code: the compatibility matrix

Act as our integration checker. Make a compatibility matrix for the team: one row for each agent and each system it connects to (Google Drive, Gmail, Google Calendar, Firecrawl, the built-in browser, our website). For each row give the interface (connector, file or browser), what the agent's job needs, the access it has now, the result of its test in Prompt B, and a status: Compatible, Conditional or Blocked — with the fix for anything not Compatible. Flag any agent with more access than its job needs. Save it as reports/compatibility-matrix.md and publish it as an artifact in plain English, coloured by status.

NOTE Least privilege is a design decision, not a setting. Ask of each connector: does this agent's job need it?

© Tertiary Infotech Academy Pte Ltd

Connect the Tools and Check Compatibility

Check your work before moving on.

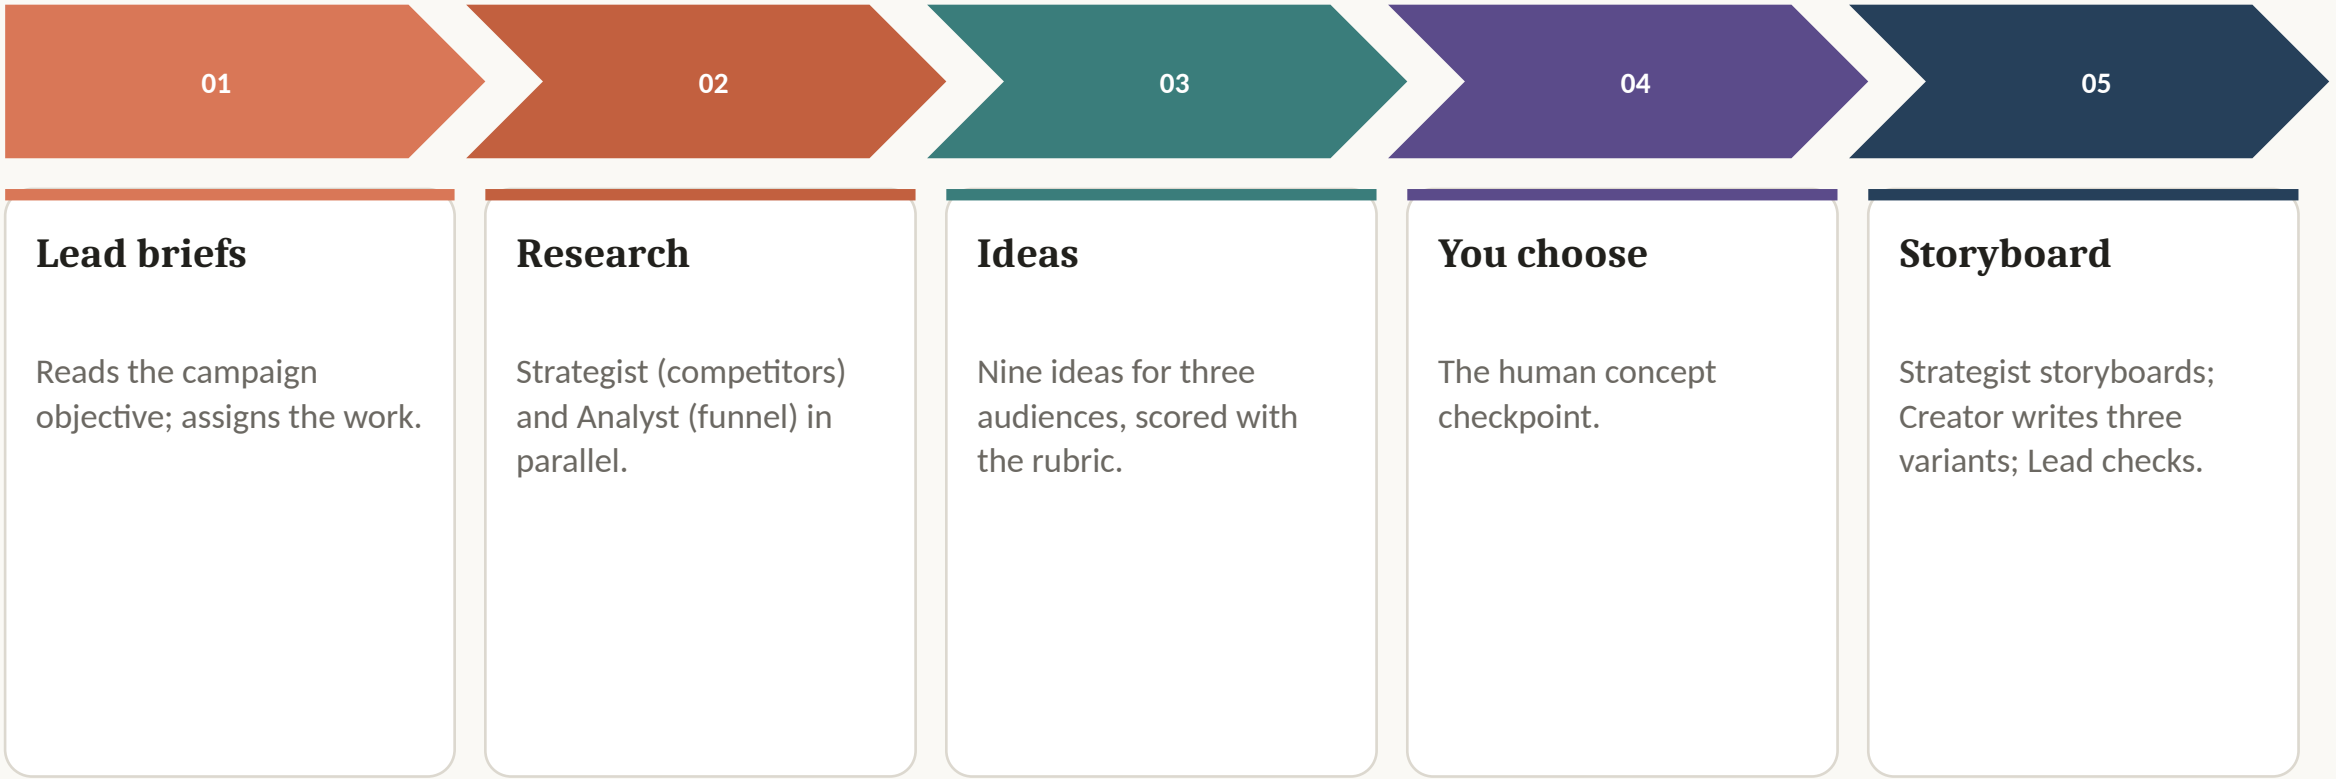
CHECK YOUR WORK

- ✓ /mcp lists Google Drive, Gmail, Google Calendar and Firecrawl; /plugin shows Marketing.
- ✓ Five checks ran, each by its own agent.
- ✓ The Gmail test is a draft in your inbox — nothing was sent.
- ✓ The analyst read the CSV from Drive, not from the local folder.
- ✓ Each agent's tools line names only the connectors its job needs.
- ✓ The compatibility matrix gives every row a status, and a fix for anything not Compatible.

© Tertiary Infotech Academy Pte Ltd

How the Team Runs a Campaign

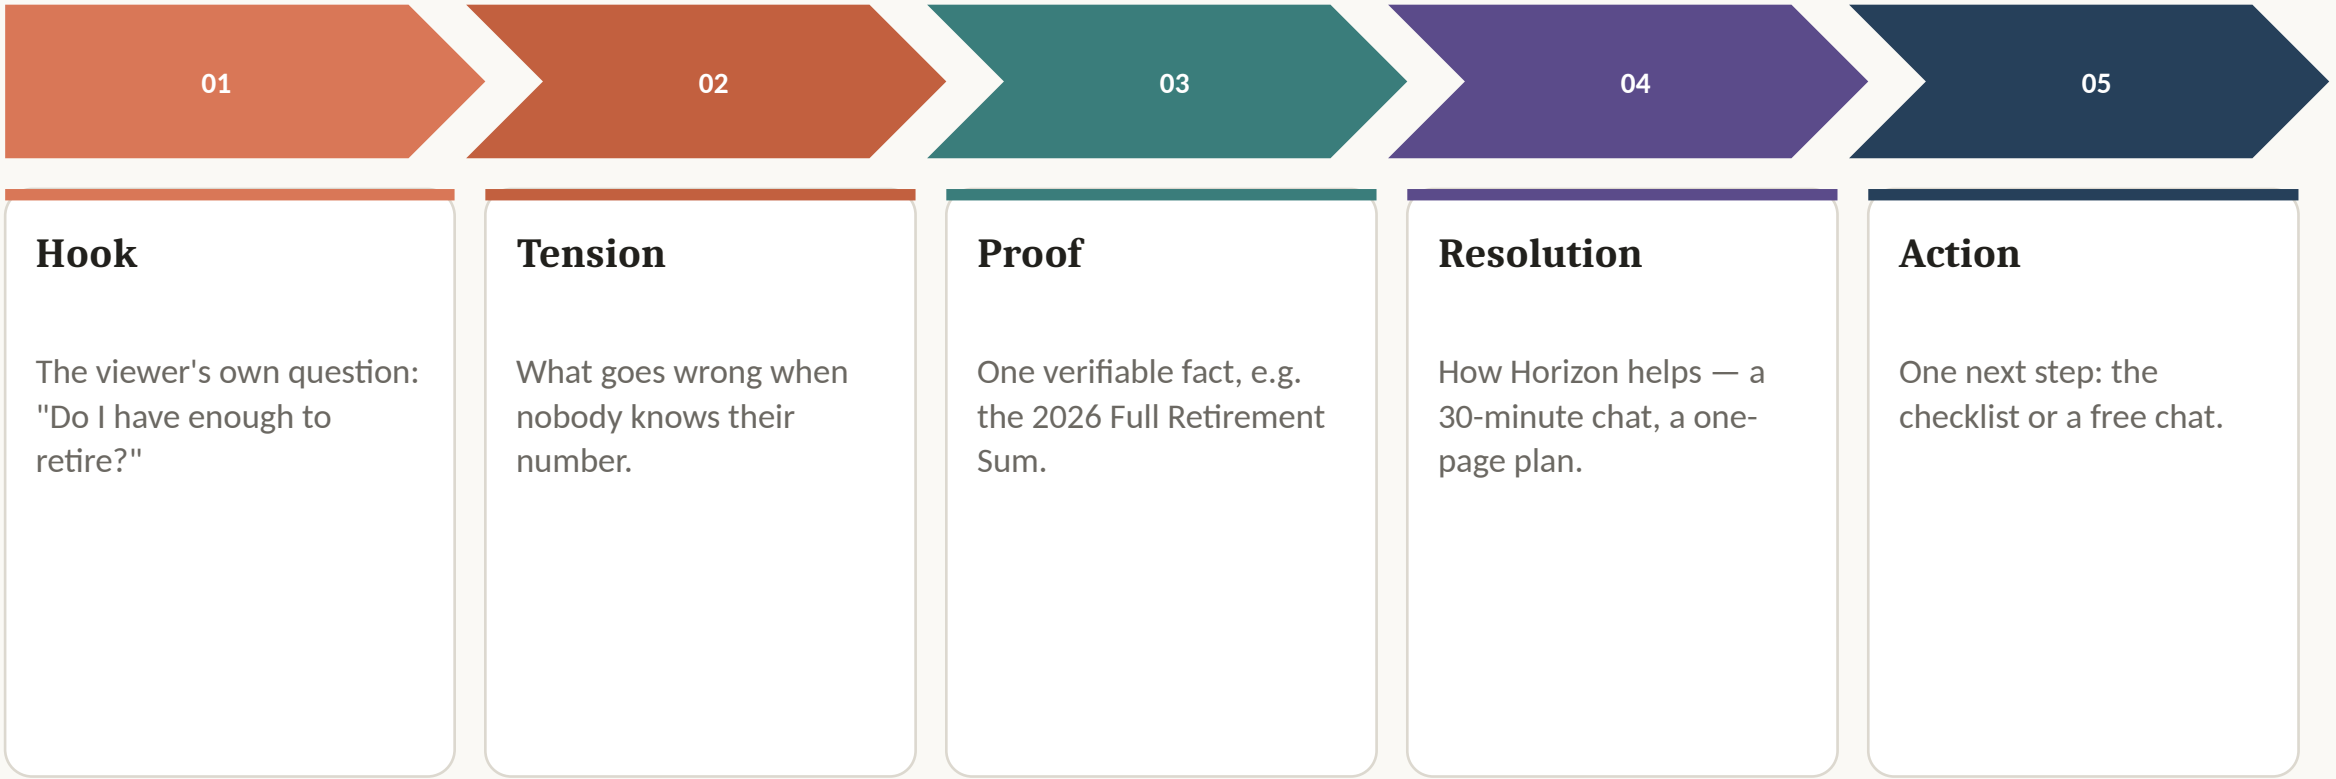
Campaign 1, "Know Your Number": research in parallel, a person chooses, the team makes it concrete.



© Tertiary Infotech Academy Pte Ltd

The Five-Beat Storyboard

Every channel tells the same story with the same claims.



© Tertiary Infotech Academy Pte Ltd

Financial Content Is Regulated

The Lead checks every piece against twelve rules (compliance-checklist.md).

RULE	WHAT IT MEANS	IF BROKEN
C1 · Not advice	General information only — say so	Critical
C2 · No promises	No guaranteed or "risk-free" returns	Critical
C3 · Label illustrations	Assumptions shown; not guaranteed	Critical
C5 · No performance	Never client returns or "average" returns	Critical
C9 · Consent	Marketing only to people who opted in	Critical
C11 · Facts sheet	Every figure from facts-2026.md	Critical

TIP MAS Guidelines on Standards of Conduct for Digital Advertising Activities apply from 25 March 2026.

© Tertiary Infotech Academy Pte Ltd

LAB 05 · HANDS-ON

Campaign 1: Research to Storyboard, with Subagents

LAB 05 | DAY 1 · ~45 MIN | USE: CLAUDE CODE

THE STORY SO FAR

The team's first campaign: 'Know Your Number', November to January. Rachel wants research, ideas for three audiences, her choice, and a storyboard every channel follows.

YOU'LL BUILD

competitor research, the funnel, nine scored ideas, the campaign brief and a five-beat storyboard with three channel variants — shown as artifacts, with a Word copy of each

Surface: Claude Code (subagents in parallel)

Features: the Lead delegates · strategist and analyst in parallel · nine scored ideas · a human concept checkpoint · a five-beat storyboard

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 5". Run all of this lab's prompts in it.
2. **Brief the Lead** The briefs are already in data/. You are the client; your session is the Marketing Team Lead.
3. **Research in parallel** Copy and paste Prompt A. Watch the strategist and the analyst run at the same time.
4. **Read the ideas** Nine ideas, scored with the rubric — the top five open as an artifact beside the chat. The Lead stops for you.
5. **Choose** Pick one idea in the artifact's Ideas tab and note its number and why — the human concept checkpoint.
6. **Storyboard** After you have chosen, copy and paste Prompt B in the chat, replacing <number> and <one line> with your choice.
7. **Decide as Rachel** After you have read the issues, copy and paste Prompt C. You confirm what is true about Horizon; the team removes or fixes the rest.

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents

PROMPT A — Claude Code: research and ideas

Act as the Marketing Team Lead for the campaign in data/campaign-objective.md. Use subagents, in parallel:
- growth-strategist: competitor-scan → save research/competitive-analysis.md.
- growth-analyst: the funnel in data/firm-metrics.csv — what 40 chats a month needs → save research/funnel.md.
When both finish, the growth-strategist proposes three ideas each for young professionals, mid-career families and pre-retirees, scored with data/idea-rubric.md and tagged with a funnel stage (content-marketing) → strategy/ideas.md. Publish the top five as an artifact — one card per idea with its score, persona and funnel stage — and a box at the top that tells me what to do next. Then stop. I will choose. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents (cont.)

PROMPT B — Claude Code: the storyboard

I choose idea <number>, because <one line>. Use subagents:
- growth-strategist: campaign-brief → strategy/campaign-brief.md, then a five-beat storyboard with data/storyboard-template.md → strategy/storyboard.md.
- content-creator: three channel variants that keep the same claims — LinkedIn, Facebook, a 60-second video. Then run fact-check and fin-compliance yourself and show me the issues before anything is final. Publish the brief and the storyboard as an artifact with a tab for the brief (goal, audience, message, call to action); a tab for the five beats side by side, each with its claim and source, and the three variants below; and the issues in a table. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents (cont.)

PROMPT C — Claude Code: decide on the issues

I confirm "no obligation" and the free checklist — add them to the facts sheet (data/facts-2026.md). Delete any other claim that is not on the facts sheet. Fix all the Medium and Low issues, then show me the updated storyboard and publish it as an artifact. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Agents diverge; a person decides. The concept checkpoint is where the campaign becomes yours.

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents (cont.)

FOLLOW-UP — type any of these in the chat

After Prompt A: Explain this result in three sentences. Which idea would you pick, and why?

Why does idea <number> score higher than idea <number>?

What does the funnel say is our biggest problem?

After Prompt B: Explain the issues in plain English. What do I need to decide?

After Prompt C: What changed since the last version? Is anything still not ready for Rachel?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Campaign 1: Research to Storyboard, with Subagents

Check your work before moving on.

CHECK YOUR WORK

- ✓ The strategist and the analyst ran in parallel.
- ✓ The competitor analysis and the funnel are done, with sources.
- ✓ Nine scored ideas, each with its author.
- ✓ You chose the idea and recorded why.
- ✓ The storyboard has five beats; every claim has a source.
- ✓ You confirmed two claims; after Prompt C no High issue is left.

© Tertiary Infotech Academy Pte Ltd

Where You Are Now

The team is integrated and has run its first campaign. Now every campaign tests the joins — and finds what breaks.

Assessed and integrated · Labs 1–4

A studio, your copy of the website, six agents with instructions, skills and connectors, and a compatibility matrix.

First integrated run · Lab 5

Research in parallel across Firecrawl and Drive, nine scored ideas, your choice, a storyboard.

What is next

Test each integration with real campaign work, read the signs of errors, troubleshoot. That is Topic 2.

© Tertiary Infotech Academy Pte Ltd

TOPIC 2

Testing Marketing Integrations and Troubleshooting Errors

1. Integration tools, techniques and protocols
2. Every campaign is an integration test
3. Personas, cadence and a blog: test the data paths
4. Social, newsletter, landing page and video: test the gate and the APIs
5. Signs of integration errors, and basic troubleshooting

© Tertiary Infotech Academy Pte Ltd

Integration Tools and Techniques

The basic toolkit you already used in Labs 1–5 — and the techniques that keep an integration safe.

Connect

Customize → Connectors and /plugin. /mcp lists what is connected; /status shows the account.

Limit

Each agent's tools line and disallowedTools — least privilege, written down.

Standardise

Skills (SKILL.md): one procedure, applied the same way by every agent.

Guard

Hooks in .claude/settings.json run before every tool call — the approval gate.

Rehearse

Dry runs: publish.mjs prints the exact request before anything is sent.

Schedule

Routines run a skill on the clock and leave a draft for a person.

© Tertiary Infotech Academy Pte Ltd

Every Campaign Is an Integration Test

Each lab joins components and checks a known answer. If the answer is wrong, an interface is wrong.

LAB	COMPONENTS JOINED	PASS WHEN
6	Survey CSVs → Growth Analyst	Counts are made with code; every claim has an n and a file
7	Benchmarks → Analyst → Strategist	A frequency per channel with evidence; no week over 12 hours
8	Strategist → Creator → website	The article shows on your site; every figure is on the facts sheet
9	Creator → Gmail; checklist → site	Three drafts in Gmail, nothing sent; the checklist is on your site
10	Team → gate → LinkedIn, Facebook	The gate stops an agent approving; your approval is stored with a hash
11	Contacts CSV → Analyst → Gmail	31 of 40 recipients — consent = yes and not unsubscribed
12	Script → render → YouTube	1920×1080, under 90 s; uploads as private

TIP Write the expected result before you run the test — then compare.

© Tertiary Infotech Academy Pte Ltd

Evidence Before Personas

Lab 6: the Growth Analyst builds personas the way an analyst would.

01

Collect

150 survey answers, 60 enquiries, 200 checklists.

02

Count

With code. Every claim has an n and a file.

03

Segment

By age band and life stage.

04

Persona

Worries, channels, formats, trust — with evidence.

05

Requirement

What each persona needs from every piece.

© Tertiary Infotech Academy Pte Ltd

Observed, Inferred or Low Confidence?

Three labels keep a persona honest.

LABEL	EXAMPLE	USE IT
Observed	30–39s chose Facebook, LinkedIn and YouTube equally (survey, n = 44)	As fact, with the count
INFERENCE	"They are short of time on weekdays" — not asked	As a hypothesis to test
LOW CONFIDENCE	The 60+ group (n = 12)	Never alone; get more data
Verbatim	"I am 54 and not sure if my CPF will be enough." (E-id)	Exactly as written, with the id

TIP Ask "which rows show that?" of anything that sounds insightful.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence

LAB 06 | DAY 1 · ~40 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Jun Wei drafted personas from instinct: 'busy young parents who love TikTok'. Rachel is not convinced. There are 150 survey answers, 60 enquiries and 200 checklists.

YOU'LL BUILD

three evidence-based personas and the content spec — shown as artifacts, with a Word copy of each

Surface: Claude Code (three analyst subagents in parallel)

Features: three analyst sub-agents, one per data file · counts, not impressions · confidence labels · content requirements

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 6". Run all of this lab's prompts in it.
2. **Check the data** The CSVs and templates are already in data/. Open one and skim it.
3. **Analyse in parallel** Copy and paste Prompt A — three analyst sub-agents, one per file.
4. **Challenge a claim** Pick one persona statement and ask: "which rows show that?"
5. **Write the guides** After the persona cards appear, copy and paste Prompt B: a writing guide for each persona — how to talk to them.
6. **Check the gaps** The 60+ group must be LOW CONFIDENCE.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence

PROMPT A — Claude Code: three analysts

Use three growth-analyst subagents in parallel, one per file: data/survey-responses.csv, data/enquiries.csv, data/checklist-results.csv. Each uses audience-insights: counts with code, n and the file on every claim, INFERENCE and LOW CONFIDENCE labels, enquiry quotes verbatim with ids. Combine the results into three personas with data/persona-template.md → strategy/personas.md. Publish the personas as an artifact: one card each, with the key counts as small bar charts and the confidence label in colour. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence (cont.)

PROMPT B — Claude Code: requirements

Use the growth-strategist and the content-creator subagents to write a writing guide for each persona (data/content-spec-template.md → strategy/content-spec.md): what they most want to know, where to reach them, what to make and how long, the tone, the proof that convinces them, what we ask them to do next, and what we never say. Back each answer with the data: which file, and how many people. Publish it as an artifact: one card per persona with a one-line summary on top and the answers under plain headings, then one table comparing the three personas side by side. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Ask "which rows show that?" of anything that sounds insightful. If the agent cannot point to rows, it is an inference.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence (cont.)

FOLLOW-UP — type any of these in the chat

Explain these personas in three sentences. Which one should we target first, and why?

Which survey answers show that <a sentence from a persona>?

Why is the 60+ group marked low confidence?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst: Personas from Evidence

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three analyst sub-agents ran, one per file.
- ✓ Every persona statement carries a count and a file.
- ✓ The 60+ group is LOW CONFIDENCE (n = 12).
- ✓ The most-unticked checklist item is the retirement number.
- ✓ Enquiry quotes are verbatim, with ids.
- ✓ The writing guide answers every question for every persona.

© Tertiary Infotech Academy Pte Ltd

Finding the Right Frequency

Lab 7: the Analyst finds where more posts stop paying; the Strategist plans four weeks.



TIP A third LinkedIn post a week brings more unfollows than value. Weekly emails make unsubscribes jump.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar

LAB 07 | DAY 1 · ~35 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Jun Wei has 12 hours a week and Rachel needs two business days to review anything. How often should each channel run, and what goes out when?

YOU'LL BUILD

how often to post on each channel, and a four-week calendar — shown as artifacts, with a Word copy of each

Surface: Claude Code (subagents)

Features: marginal value vs fatigue · team capacity · review buffers · a four-week calendar

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 7". Run all of this lab's prompts in it.
2. **Check the data** The assets are already in data/. Open one and skim it.
3. **Find the sweet spot** Copy and paste Prompt A.
4. **Question it** Why is the newsletter not weekly? The answer is in the unsubscribes.
5. **Plan four weeks** After the analyst recommends a cadence, copy and paste Prompt B.
6. **Check the buffers** Every review date is two business days ahead; nothing on a holiday.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar

PROMPT A — Claude Code: the analyst

Use the growth-analyst subagent. From data/channel-benchmarks.csv and data/team-capacity.md, recommend how often to publish on LinkedIn, Facebook, the newsletter and YouTube. Compare the extra clicks from one more post, the unfollows or unsubscribes it costs, and the hours → strategy/cadence.md. Publish it as an artifact: a chart per channel of the extra clicks against the unfollows, and the hours. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar (cont.)

PROMPT B — Claude Code: the strategist

Use the growth-strategist subagent to build a four-week calendar from Monday 2 November 2026 in strategy/calendar.csv (columns as data/calendar-format.csv), from the cadence, strategy/storyboard.md and strategy/content-spec.md. review_by two business days before each date; no public holidays; no week over the team's hours. Publish the four weeks as an artifact: a calendar grid by channel, review dates marked, hours per week. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Frequency is a capacity decision as much as an audience one. A calendar the team cannot deliver is fiction.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar (cont.)

FOLLOW-UP — type any of these in the chat

Why is the newsletter not weekly? Show me the numbers.

What happens if we add a third LinkedIn post a week?

Which week is busiest, and does the team have the hours?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Growth Analyst and Strategist: Cadence and Calendar

Check your work before moving on.

CHECK YOUR WORK

- [x] The cadence gives a frequency per channel with the evidence.
- [x] A third LinkedIn post a week costs far more unfollows than it earns.
- [x] It explains the unsubscribe jump at weekly emails.
- [x] The calendar covers four weeks in the required columns.
- [x] Every review_by date is two business days before publishing.
- [x] No week is over 12 hours.

© Tertiary Infotech Academy Pte Ltd

From Storyboard to Content

Lab 8: the first real piece — a blog article every post can link to.

One prompt, the whole job

Research, draft, fact-check, fix and post — a quick blog in one go. A person still reads it last.

Straight onto your copy of the site

The article lands on the copy of Horizon's site in your folder, opened in your browser. The public site never changes.

Your skills know Horizon

blog-post, sunny-voice, fact-check and fin-compliance: the voice, the facts and the rules.

© Tertiary Infotech Academy Pte Ltd

Content Creator: A Quick Blog Article

LAB 08 | DAY 1 · ~40 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Jun Wei has the storyboard and the personas. Rachel wants the first real piece today: a blog article that people searching for CPF top-ups will find — and that every post can link to.

YOU'LL BUILD

a 500-700-word blog article with dated sources, checked and fixed and posted on your own copy of Horizon's website, opened in your browser, with a Word copy

Surface: Claude Code (Firecrawl, your Horizon skills) → your copy of Horizon's website

Features: one prompt runs the whole job: research · draft · fact-check and fin-compliance · fix · post to the blog on your copy of the website · a person reads it last

© Tertiary Infotech Academy Pte Ltd

Content Creator: A Quick Blog Article

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 8". Run all of this lab's prompts in it.

2. **Write and post it** Copy and paste the prompt. Claude researches, writes, checks and fixes the article, then posts it to your copy of Horizon's website.

3. **Read it on the site** Your copy of the website opens in your browser: your article is first under "From our blog". Say what to change, if anything.

© Tertiary Infotech Academy Pte Ltd

Content Creator: A Quick Blog Article

PROMPT — Claude Code: a quick blog

Write a quick blog article for Horizon and post it to the website. Use the blog-post, sunny-voice, fact-check and fin-compliance skills.
1. Pick one topic from strategy/storyboard.md that a Singaporean would search for. Use Firecrawl to find two current sources (CPF Board, MAS or MoneySense), each with its date.
2. Write 500-700 words: a title under 60 characters, every figure from data/facts-2026.md, one call to action (book a free 30-minute chat) and the full disclaimer → content/blog/<slug>.md.
3. Run fact-check and fin-compliance; fix every issue and list each fix in one line.
4. Add it to the blog in web/site/index.html as one new post, built like the posts already there. Change nothing else on the site. Then start my website if it is not running, and open it. Write for a business owner, in plain English. Save a Word copy (.docx) of the article.

NOTE One prompt can run the whole job — research, writing and checks. A person still reads the result before it goes anywhere.

© Tertiary Infotech Academy Pte Ltd

Content Creator: A Quick Blog Article (cont.)

FOLLOW-UP — type any of these in the chat

Summarise the article in three sentences. Who is it for?

Why did you pick this topic, and who will search for it?

What did the checks find, and how did you fix it?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Content Creator: A Quick Blog Article

Check your work before moving on.

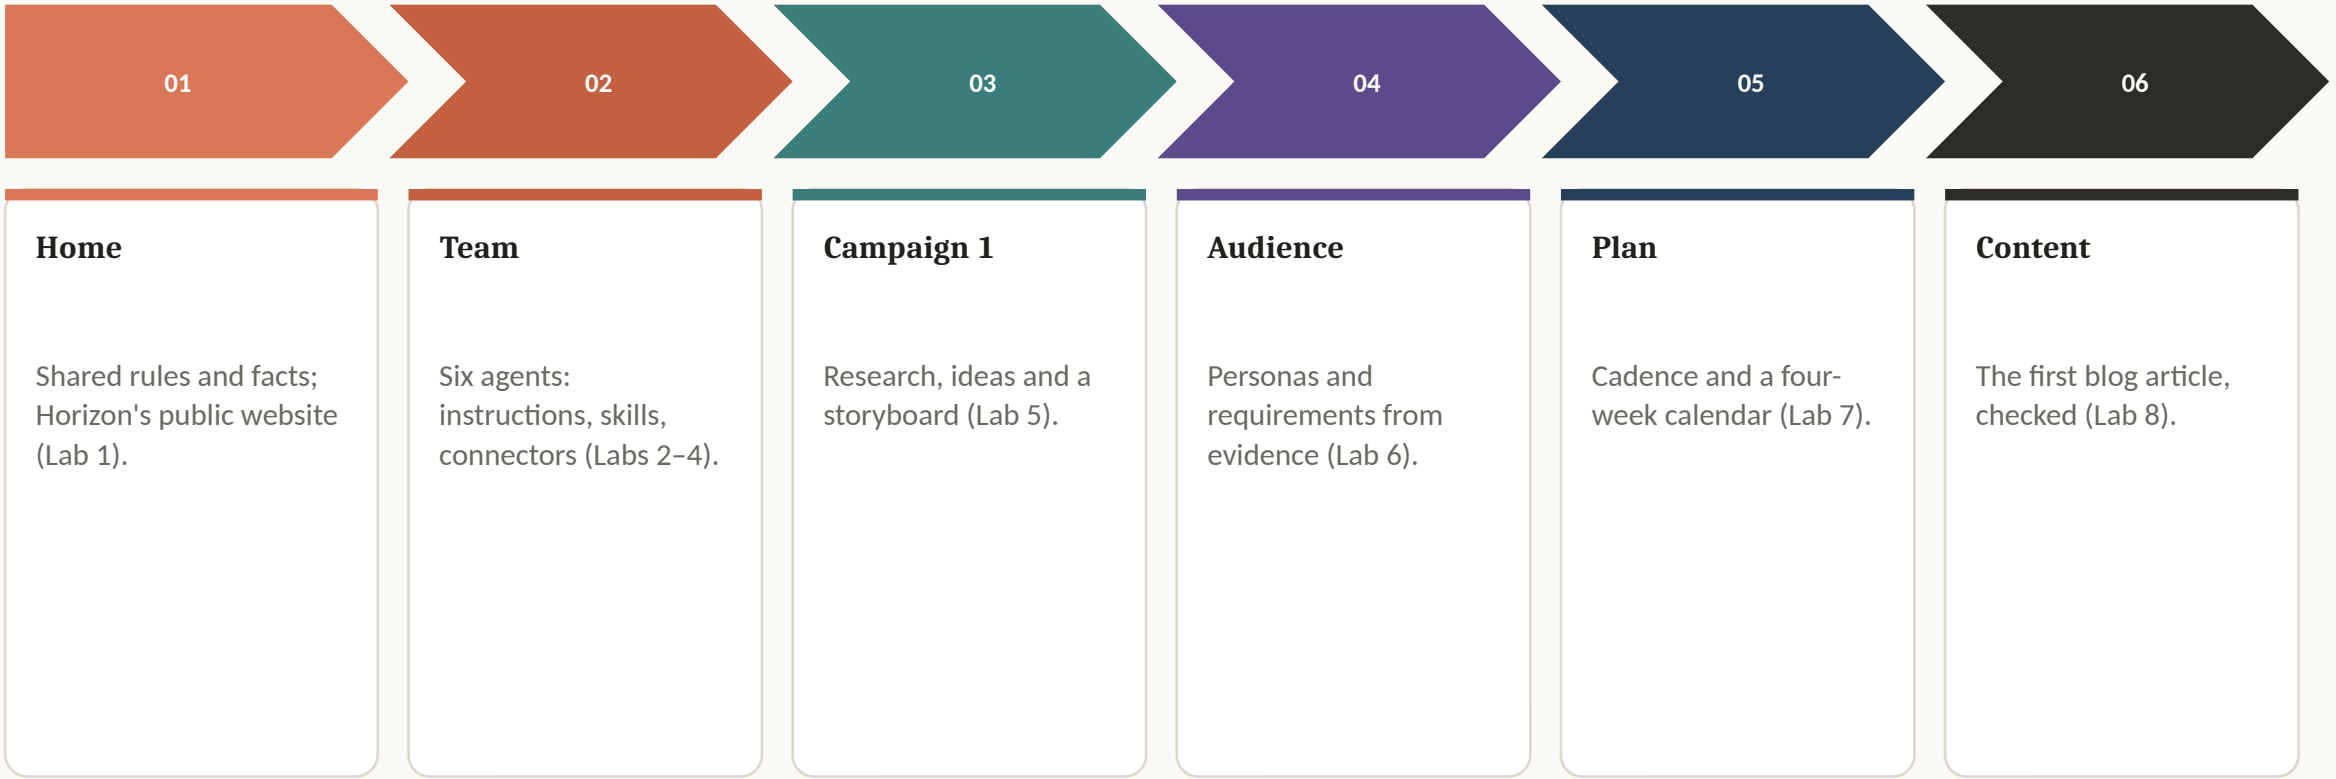
CHECK YOUR WORK

- ✓ Two sources, each with a date and a link.
- ✓ The title is under 60 characters.
- ✓ Every figure is from Horizon's 2026 facts.
- ✓ The fixes are listed; no Critical or High issue is left.
- ✓ One call to action and the full disclaimer at the end.
- ✓ Your article is first under "From our blog" on your copy of the website.

© Tertiary Infotech Academy Pte Ltd

Day 1: The Team Is Ready

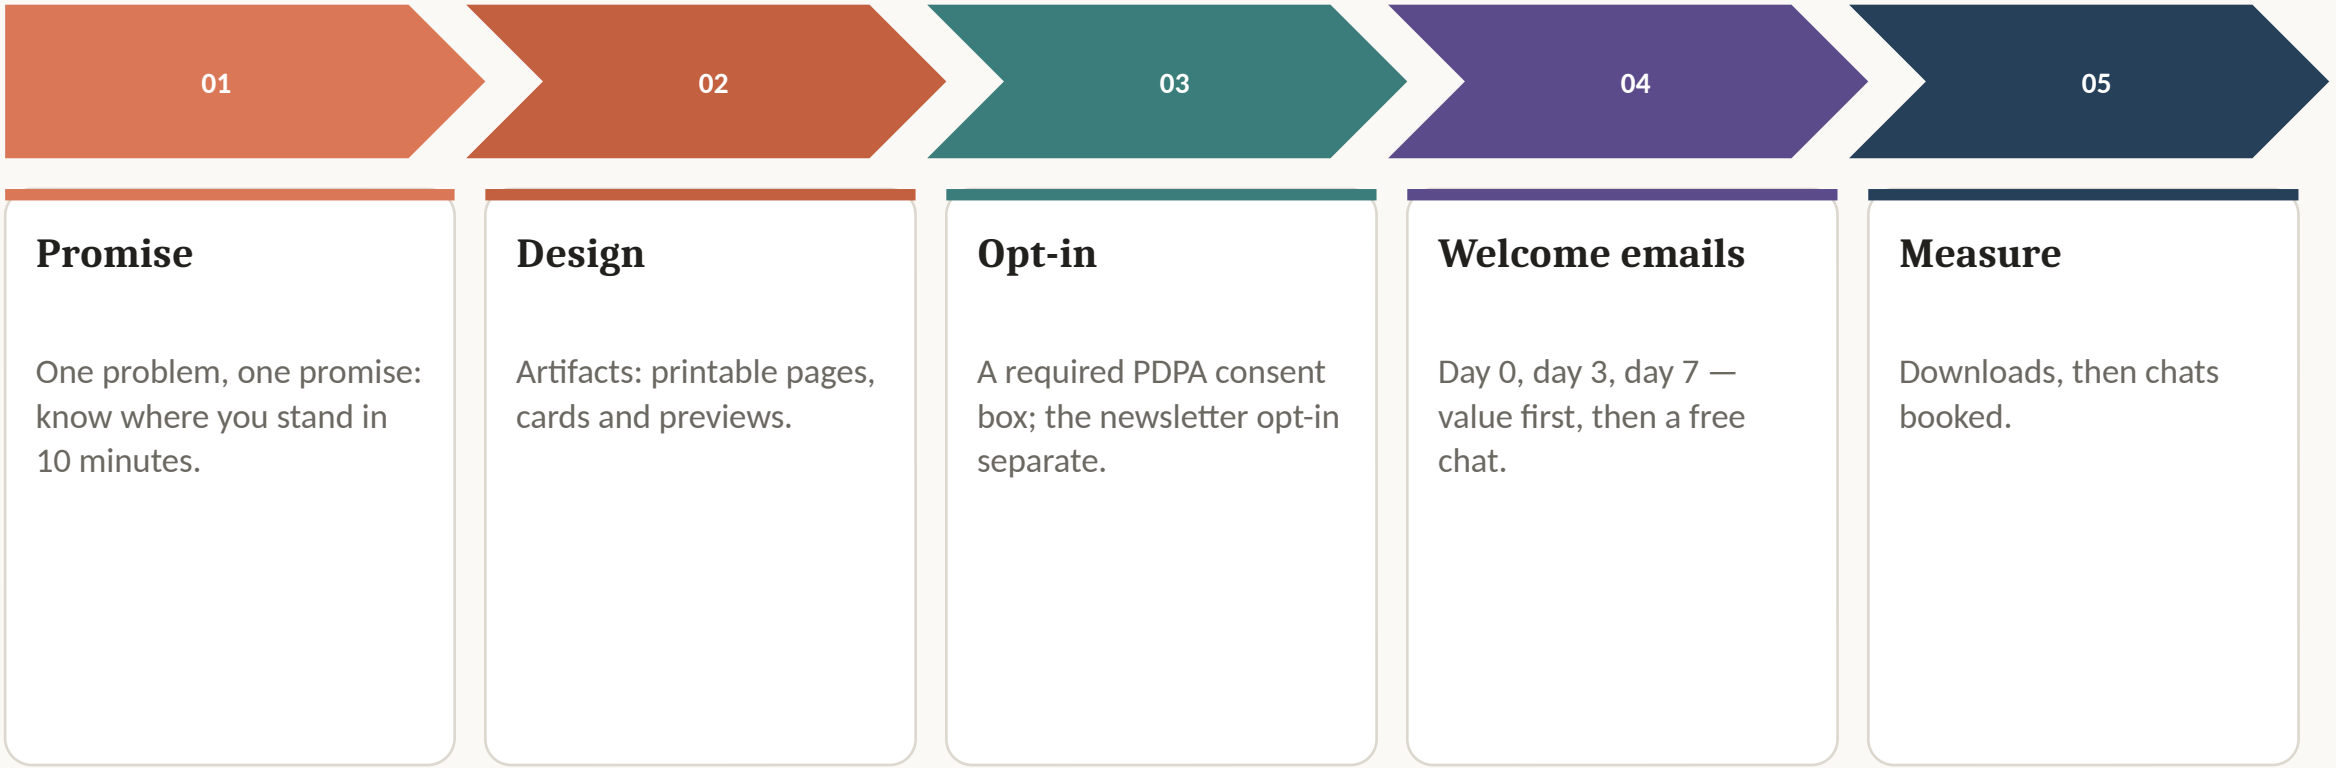
Everything the campaigns need tomorrow is in place.



© Tertiary Infotech Academy Pte Ltd

Content That Earns a Sign-up

Lab 9: the Money Check-up Checklist and the emails that follow it.



TIP Lab 11 builds the landing page that offers it.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails

LAB 09 | DAY 2 · ~35 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Day 2. Before the campaign weeks start, Rachel wants the thing that turns readers into leads: the Money Check-up Checklist, designed to be downloaded, and the emails that follow it.

YOU'LL BUILD

the Money Check-up Checklist as a printable A4 page (artifact and PDF) and on your copy of Horizon's website, and three welcome emails — shown as artifacts, with a Word copy of each

Surface: Claude Code (artifact, Gmail drafts, Marketing plugin, your Horizon skills)

Features: lead-magnet skill · a printable A4 checklist as an artifact · the Marketing plugin's email sequence · newsletter and sunny-voice · drafts only

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 9". Run all of this lab's prompts in it.

2. **Design it** Copy and paste Prompt A. The checklist opens as an artifact; ask in the chat for any change.

3. **Write the welcome emails** After the checklist looks right, copy and paste Prompt B. The three emails also appear as drafts in your Gmail.

4. **Check it** After the three emails are saved, copy and paste Prompt C.

5. **Put it on the website** After the checks pass, copy and paste Prompt D. Your copy of the website opens in your browser, with the checklist behind an email box.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails

PROMPT A — Claude Code: the checklist

Use the lead-magnet skill. Turn data/checklist-items.md into the "Money Check-up Checklist": a title with one promise, the ten items as tick boxes, one figure from data/facts-2026.md, the next step (book a free chat) and the full disclaimer. Lay it out as a printable A4 page in the colours and fonts in data/brand.md and publish it as an artifact. Save a PDF of it in content/lead-magnet/. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

PROMPT B — Claude Code: the welcome emails

Use the Marketing plugin's email sequence to write three welcome emails for people who download the checklist: day 0 (here it is), day 3 (the item most people miss — the retirement number), day 7 (book a free chat). Follow the newsletter and sunny-voice skills → content/email/welcome-sequence.md. Then, using the Gmail connector, create the three emails as DRAFTS in my Gmail, addressed to me. Send nothing. Publish the three emails as an artifact: one tab per email, laid out as it will look in the inbox. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

PROMPT C — Claude Code: check it

Run fact-check and fin-compliance on the checklist text and the three emails. Show the issues in one table, with a fix for each. Once I agree, make the fixes in the files and in the Gmail drafts. Publish the table as an artifact, coloured by severity. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

PROMPT D — Claude Code: put it on the website

Add the checked Money Check-up Checklist to Horizon's website in web/site/index.html, as a new section before the reviews, in the site's style:
the title and promise, what's inside, and an email box with a required consent tick box. Once a valid email is entered, show the ten items as tick boxes and a "Print or save as PDF" button. Add "Free checklist" to the menu. Change nothing else on the site. Then start my website if it is not running, and open it.

NOTE A lead magnet earns an email address. Give real value first; the call to action comes last.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

FOLLOW-UP — type any of these in the chat

Which checklist item do people most often miss, and why does it matter?

Read the day 3 email as a 55-year-old. Is anything unclear?

Explain each issue in one line: what is wrong, and the fix.

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

The Lead Magnet: a Checklist and Its Welcome Emails

Check your work before moving on.

CHECK YOUR WORK

- ✓ The checklist uses Horizon's colours and all ten items.
- ✓ It makes one promise and ends with the full disclaimer.
- ✓ Three emails, each with a subject under 50 characters, are drafts in your Gmail — nothing was sent.
- ✓ Each email has one call to action, unsubscribe and the disclaimer.
- ✓ fin-compliance shows no Critical or High issue.
- ✓ The checklist is on your copy of the website: email and consent first, then the ten tick boxes.

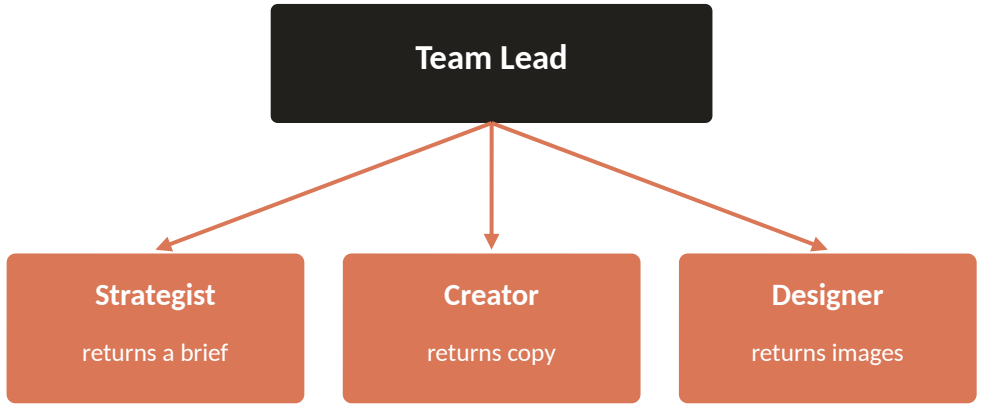
© Tertiary Infotech Academy Pte Ltd

Subagents vs Agent Teams

Both run agents in parallel, each with its own context. The difference is who they talk to.

SUBAGENTS

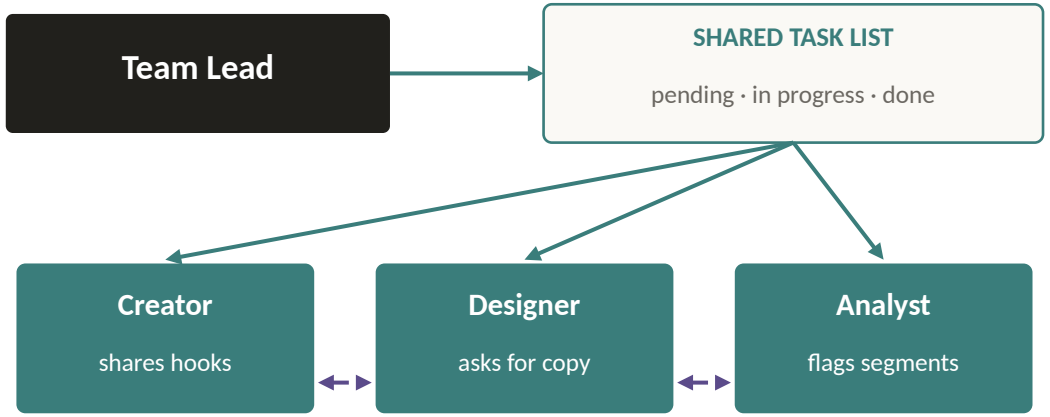
Report back to one main agent



Each works in its own context and returns only its result. They do not talk to each other.

AGENT TEAM

A lead, a shared task list, teammates who talk



Teammates claim tasks and message each other directly — e.g. a reviewer sends fixes straight to the writer.

Use subagents when only the result matters. Use a team when the agents need to talk.

© Tertiary Infotech Academy Pte Ltd

Which Pattern When?

From the Claude Code documentation, applied to the team.

	SUBAGENTS	AGENT TEAM
Talks to	Returns a result to the Lead	Teammates message each other directly
Coordination	The Lead manages everything	A shared task list; teammates claim tasks
Best for	Research, analysis, one-off reviews	Copy ↔ design ↔ review loops
Cost	Lower: results are summarised back	Higher: every teammate is its own Claude session
Campaigns	Lab 5 (and sub-agents in Claude: Labs 6, 15)	Labs 10, 11, 12

TIP Agent teams are experimental; your studio folder already has them switched on.

© Tertiary Infotech Academy Pte Ltd

Running an Agent Team

In Claude Code. Describe the team in plain words.

01

Already on

Your studio folder switches agent teams on — nothing to set.

02

Name the agent types

"Spawn teammates using the agent types content-creator and creative-designer."

03

Watch

↑ ↓ and Enter open a teammate; Ctrl+T shows the task list.

04

Steer

Message any teammate directly; tell the Lead to wait.

05

Shut down

"Ask the teammates to shut down" when done.

TIP Teammates load CLAUDE.md, skills and connectors — but not the Lead's conversation. One team per session.

© Tertiary Infotech Academy Pte Ltd

Human in the Loop, Enforced

A rule in a prompt persuades. The publisher kit builds a gate that holds.

Only people approve

You type "approved <id> by <your name>". The gate reads your own message — an agent cannot write one.

The approved words go out

The approval stores a hash. Edit one word and publish --live refuses until a person approves again.

A hook stops agents

The gate blocks agents from approving, editing the approval records, or publishing unapproved items.

© Tertiary Infotech Academy Pte Ltd

How the Gate Works

Already in your studio folder. You never run it — you just type.

01

Submit

The Lead puts a checked piece up for review.

02

You read

The post, the image, the email or the video.

03

You type

"approved <id> by <your name>" — or "changes <id>: what to fix".

04

Recorded

Your name, the time and a fingerprint of the exact words.

05

Publish

Only approved, unchanged pieces. Edit one word and it needs you again.

TIP The gate is a hook: it reads your own message, and blocks any agent that tries to approve or publish on its own.

© Tertiary Infotech Academy Pte Ltd

The Publishing APIs

publish.mjs is a dry run by default. Tokens live only in .env.

CHANNEL	ENDPOINT	PERMISSION
LinkedIn	POST api.linkedin.com/rest/posts (LinkedIn-Version: YYYYMM)	w_member_social — your own profile
Facebook	POST graph.facebook.com/<version>/<page-id>/feed	Page token: pages_manage_posts
YouTube	Resumable upload to youtube/v3/videos, privacy private	youtube.upload; unaudited projects stay private
Newsletter	No API — HTML + text files, then Gmail drafts	Consent = yes, unsubscribed = no

TIP No accounts? Every lab works in dry-run mode: you see the exact request that would be sent.

© Tertiary Infotech Academy Pte Ltd

Protocols at Each Interface

Two components integrate only when both accept the same protocol at the interface between them.

INTERFACE	PROTOCOL	ACCEPTED WHEN
Claude ↔ connectors	MCP (Model Context Protocol), JSON-RPC	Listed in /mcp; the agent's tools line names it
You ↔ Google	OAuth 2.0 sign-in with scopes	Your own account; the scope covers the job
Publisher ↔ platforms	HTTPS REST + JSON, bearer token	Right version header; token in .env only
Claude ↔ the gate	Hook: JSON on stdin, exit code	Exit 2 blocks the call and says why
Agents ↔ files	Markdown + YAML frontmatter; CSV UTF-8	Required fields present; columns match
Site ↔ browser	HTTP on localhost:8080	Page opens; links carry UTM tags

TIP Protocol mismatches show up as errors at the interface — not in the model's answer.

© Tertiary Infotech Academy Pte Ltd

Signs of Integration Errors

Most failures announce themselves. Read the sign, then go to the interface it points at.

SIGN	LIKELY CAUSE	WHERE TO LOOK
"I have no Drive tool"	Connector not on the tools line, or an old session	Agent file; /mcp; new session
"Insufficient scope" / 403	Signed in without the access the job needs	Reconnect; grant the scope
A figure not on the facts sheet	The agent used memory, not data/	fact-check; the source file
40 recipients, not 31	The consent filter was not applied	The contacts CSV columns
Publish refused: hash mismatch	The words changed after approval	approvals.csv; approve again
Garbled S$ or dates	Wrong encoding or date format	Open the CSV; re-save as UTF-8

TIP A refusal from the gate is the integration working — not a bug.

© Tertiary Infotech Academy Pte Ltd

Campaign 2: A Social Media Week

Lab 10, a demo: one LinkedIn and one Facebook post, each with a visual.

01	02	03	04	05
Lead	**Content Creator**	**Creative Designer**	**Lead review**	**You**
Reads the brief and the calendar; assigns the posts.	Two posts with hooks, UTM links, disclaimers.	One image per post, with alt text.	fact-check and fin-compliance; fixes go back.	Approve on the review page and paste the line in the chat; then post it yourself (or, optionally, through the API).

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo)

LAB 10 DAY 2 · ~40 MIN USE: CLAUDE CODE

THE STORY SO FAR

Week 1 of the campaign starts Monday. As a demo, one LinkedIn post and one Facebook post, each with a visual — and Rachel's one rule: nothing goes out unless a person approved that exact version.

YOU'LL BUILD

two posts with images, reviewed, approved by you and posted by you (or, optionally, through the API) — shown as artifacts, with a Word copy of each

Surface: Claude Code (agent team) → LinkedIn, Facebook

Features: the publisher kit and its hook · an agent team: creator, designer, Lead · teammates message each other · human approval → you post

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo)

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 10". Run all of this lab's prompts in it.
2. **Test the gate** Ask Claude to approve a post itself. The gate (set up in your folder since Lab 1) must block it.
3. **Start the team** Copy and paste Prompt A. The Lead starts two teammates and tells you as each one starts and finishes.
4. **Watch them talk** The review page opens with a Team chat: the designer asking the creator for each hook, and the Lead's fixes.
5. **Approve as a person** Review in the artifact or the review email. Click Approve or Request changes, then Copy and paste into the chat.
6. **Get ready to post** After you have approved the posts, copy and paste Prompt B. Each post gets a Copy button, and the folder with the two images opens.
7. **Post it yourself** Paste the text, add the image, Post — on LinkedIn and a test Facebook Page. Or stop at the preview.
8. **Optional: post through the apps** Follow "Optional: post through LinkedIn and Facebook" in this lab's README.

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo)

PROMPT A — Claude Code: the social team

Create an agent team for Week 1 social from data/social-brief.md and strategy/calendar.csv. Spawn two teammates using the agent types content-creator and creative-designer; you are the Lead.
- content-creator: 1 LinkedIn + 1 Facebook post with linkedin-post, facebook-post and copywriting, one file each in content/social/week-01/, with the frontmatter in data/publishing-spec.md. Each post links to my website, http://localhost:8080/, with UTM tags.
- creative-designer: one image per post (a rendered card in Horizon colours: LinkedIn 1200 x 627, Facebook 1080 x 1080) with alt text; message the creator for each post's hook before designing. Save each as a PNG beside its post and fill in the post's image and alt lines.
- You: fact-check and fin-compliance on every post and image; message fixes to the teammate who owns it. Before you start, ask me how I will review: in the artifact, or by email (then ask for the approver's address). When a post passes, submit it for review with scripts/submit.mjs. Publish the two posts as an artifact. At the top, a Team chat: the messages the teammates and you sent each other, in order, each with who sent it. Then each post as it will look in the feed, with its image and alt text; under each post put two working buttons (JavaScript): Approve turns the card green; Request changes opens a box for my note. A bar at the bottom asks my name once, builds the exact lines to paste in the chat ("approved <id> by <name>", "changes <id>: <note>") and has a Copy button that says "Copied - now paste it in the chat". Keep the lines on screen in case copying is blocked. If I chose email, send ONE review email to the approver with Gmail: the artifact link and what passed the checks. Send nothing else; publish nothing. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo) (cont.)

PROMPT B — Claude Code: ready to post

For each Week 1 post approved in review/approvals.csv, publish a "ready to post" artifact: the post text with a Copy button, its image, and the steps to post it by hand on LinkedIn or our Facebook Page. Then open the folder with the two images on my computer. Publish nothing yourself. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Subagents report to one boss; teammates also talk to each other. A designer who can ask the writer for the hook makes a better image.

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo) (cont.)

FOLLOW-UP — type any of these in the chat

Which post do you expect to do best, and why?

What did the Lead send back to the writer, and why?

Is anything here not ready to post? Explain in plain English.

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Campaign 2: Social Media Posts (Demo)

Check your work before moving on.

CHECK YOUR WORK

- ✓ The gate stopped Claude from approving.
- ✓ Two teammates ran; the Lead said when each started and finished.
- ✓ The Team chat shows the designer and the creator messaging each other.
- ✓ Two posts with images from the designer, alt text and UTM links.
- ✓ You reviewed in the artifact or by email and pasted the approval in the chat; the rows show your name and a hash.
- ✓ You posted the approved text and image yourself (or, optionally, through the API after you said go).

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter + Landing Page

Lab 11: three specialists, one email and one page, consent first.

01

Content Creator

The issue and the landing-page copy.

02

Website Designer

The checklist page, page-qa, then a Claude artifact.

03

Growth Analyst

Recipients by consent: 31 of 40.

04

Lead

Checks everything; the build waits for you.

05

Drafts

Gmail drafts — one to you to test. Nothing sent.

© Tertiary Infotech Academy Pte Ltd

Consent Decides Who Gets Mail

The PDPA, applied by an agent that counts.

Opt-in only

consent = yes and unsubscribed = no. Everyone else goes to excluded.csv with a reason.

Every email carries

An unsubscribe link, Horizon's address and the short disclaimer.

The page collects little

Name, email, a required consent box, a separate unticked newsletter opt-in.

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page

LAB 11 DAY 2 · ~40 MIN USE: CLAUDE CODE

THE STORY SO FAR

The Sunny Sunday email books chats most cheaply. November's issue promotes a new landing page for the free checklist — and must reach only the 31 people who opted in.

YOU'LL BUILD

the November newsletter, its recipient list, the landing page published, and Gmail drafts — shown as artifacts, with a Word copy of each

Surface: Claude Code (agent team) → Claude artifact, Gmail

Features: creator writes, website designer builds the lead-magnet page, analyst filters by consent, the Lead checks · drafts only

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 11". Run all of this lab's prompts in it.
2. **Start the campaign** Copy and paste Prompt A.
3. **Approve the words** Review the issue in the artifact (or the review email's link). Click Approve or Request changes, then Copy and paste into the chat.
4. **Publish the page** The website designer runs page-qa and publishes web/checklist/ as a Claude artifact.
5. **Check the list** Only people with consent = yes and unsubscribed = no are on the recipient list.
6. **Create the drafts** After the page is published and the list checked, copy and paste Prompt B.
7. **Send one test** Send only the draft addressed to yourself, and read it on your phone.

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page

PROMPT A — Claude Code: the newsletter team

Create an agent team for the November Sunny Sunday email (data/newsletter-brief.md, data/newsletter-spec.md) and its landing page (data/landing-page-brief.md). Spawn three teammates:
content-creator, website-designer and growth-analyst.
- content-creator: issue.md (id nl-2026-11) with newsletter, and the landing-page copy.
- website-designer: web/checklist/index.html with lead-magnet, web-design and landing-page from the approved copy; page-qa.
- growth-analyst: recipients.csv and excluded.csv from data/subscribers.csv. Before you start, ask me how I will review: in the artifact, or by email (then ask for the approver's address). You check everything, then submit nl-2026-11 and publish the issue as an artifact, laid out as the email, with the counts of recipients and excluded, and two working buttons (JavaScript): Approve turns the page green; Request changes opens a box for my note. A bar at the bottom asks my name once, builds the exact line to paste in the chat ("approved <id> by <name>" or "changes <id>: <note>") and has a Copy button that says "Copied - now paste it in the chat". Keep the line on screen in case copying is blocked. If I chose email, send ONE review email to the approver with Gmail: the artifact link and what passed the checks. Only after I approve, build newsletter.html and .txt and publish the page as an artifact. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page (cont.)

PROMPT B — Claude Code: Gmail drafts

Using the Gmail connector, create a DRAFT of newsletter.html for each address in recipients.csv, plus one draft addressed to me only. Send nothing. Publish a report as an artifact: how many drafts, how many people were excluded, and why, in a table.

NOTE Consent decides who gets mail — not the agent and not the deadline. Marketing to people who did not opt in breaks the PDPA.

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page (cont.)

FOLLOW-UP — type any of these in the chat

Explain the email in three sentences. What do we want readers to do?

Who was left off the mailing list, and why?

What did the landing-page check look at, and did it pass?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Campaign 3: Newsletter and a Landing Page

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three specialists worked on the campaign; the Lead checked all.
- ✓ You reviewed nl-2026-11 in the artifact or by email and approved it in the chat before the build.
- ✓ The landing page passed page-qa and is published.
- ✓ The recipient list holds 31 people; 9 are excluded, with reasons.
- ✓ The email has unsubscribe, address and disclaimer.
- ✓ Drafts only — nothing was sent to the list.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer

Lab 12: the team writes and designs; you approve twice.

Strategist and Creator

The angle, then 5–7 scenes with words on screen and a description with chapters and an AI disclosure.

Creative Designer

Storyboard, thumbnail, then the render with video-render — ffmpeg, 1920×1080, under 90 s.

Upload privately

Approve the words, then the finished video. It uploads as private; a person makes it public.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer

LAB 12 | DAY 2 · ~30 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Pre-retirees say they would rather watch than read. Rachel wants the first 'Money in Plain English' video, uploaded privately until she has watched it.

YOU'LL BUILD

the script, storyboard and thumbnail, the video, and a private YouTube upload (or dry run) — shown as artifacts, with a Word copy of each

Surface: Claude Code (agent team, render, upload)

Features: strategist angle · creator script · designer storyboard, thumbnail and render · approve twice · private upload

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 12". Run all of this lab's prompts in it.
2. **Script it as a team** Copy and paste Prompt A.
3. **Approve the script** Read the scenes, title and description in the storyboard artifact; type: approved yt-ep01-script by <your name>.
4. **Render it** After you approve the script, copy and paste Prompt B. The designer runs video-render.
5. **Approve the video** Ask the Lead to submit the video, watch ep01.mp4 to the end, then approve yt-ep01.
6. **Upload** After you approve the video, copy and paste Prompt C. It goes up as private.
7. **Disclose** In YouTube Studio, answer the altered or synthetic content question before making it public.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer

PROMPT A — Claude Code: the video team

Create an agent team for YouTube episode 1 from data/youtube-brief.md. Spawn three teammates: growth-strategist (the angle and hook), content-creator (content/video/ep01/scenes.json, 5-7 scenes, and ep01.md with title, tags, description, chapters and the AI disclosure) and creative-designer (a storyboard of the scenes and a thumbnail, using brand-visuals). You fact-check and fin-compliance the script, then submit it as yt-ep01-script. Publish nothing. Publish the storyboard as an artifact: one frame per scene with its words, visual and timing, and the thumbnail. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer (cont.)

PROMPTS B and C — Claude Code: render and upload

PROMPT B

Use the creative-designer subagent with video-render on the approved scenes.json → content/video/ep01/ep01.mp4 (data/video-spec.md). Check its length and size. Do not change approved words. Publish the checks as an artifact: length, size, resolution and every scene, PASS or FAIL.

PROMPT C

Do a dry run of the YouTube upload for yt-ep01 (scripts/publish.mjs). If I say "go", upload it for real — as PRIVATE. Publish the request, then the YouTube Studio link, as an artifact.

NOTE Approve twice: once for the words, once for the finished video. Rendering can break what the script got right.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer (cont.)

FOLLOW-UP — type any of these in the chat

Walk me through the video scene by scene, in plain English.

Which figure on screen should I double-check, and against what?

What does the AI disclosure say, and why do we need it?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Campaign 4: The YouTube Explainer

Check your work before moving on.

CHECK YOUR WORK

- [x] The strategist, creator and designer each delivered their part.
- [x] You approved the script, then the finished video.
- [x] ep01.mp4 is 1920x1080, under 90 seconds, every scene present.
- [x] Every figure on screen matches Horizon's 2026 facts.
- [x] The description has chapters, link, disclaimer and AI disclosure.
- [x] The upload is private (or the dry run printed the request).

© Tertiary Infotech Academy Pte Ltd

Where You Are Now

Seven labs, each an integration test: every piece checked by the Lead and approved by a person. But how do you change the team safely?

Tested · Labs 6–9

Data paths: personas, cadence, the blog on your site, Gmail drafts and the checklist page.

Troubleshot · Labs 10–12

The gate, the consent filter, the render and the APIs — and the signs when they break.

What is next

Fix and improve without breaking what works: methods, versions, change proposals. That is Topic 3.

© Tertiary Infotech Academy Pte Ltd

TOPIC 3

Versioning, Change Proposals and Integration Best Practices

1. Troubleshooting methodologies
2. Always on: from skill to schedule
3. Responsible AI; measure what matters
4. Version the team; propose a change; roll back
5. Agent teams at work: argue a strategy, then a QA loop

© Tertiary Infotech Academy Pte Ltd

A Troubleshooting Method

Guessing changes three things and fixes none. A method changes one thing and proves it.

01

Reproduce

Make it fail again on purpose, with the same input.

02

Isolate

One component at a time: /mcp, /status, one small test per agent.

03

Inspect the interface

Format, protocol, account and scope — where the two parts meet.

04

Fix one thing

The smallest change that could explain the fault.

05

Retest and record

Run the same test; log the cause and the fix in the playbook.

TIP In Lab 15 the QA Auditor runs this loop: fail, fix, audit again — three rounds at most.

© Tertiary Infotech Academy Pte Ltd

Five Ways to Find the Fault

Classic integration troubleshooting methods, applied to Horizon's team.

Top-down

Start from the symptom — a wrong post — and trace back: Lead, agent, skill, data.

Bottom-up

Test the lowest layer first: is the connector signed in? Does the file open? Then build up.

Divide and conquer

Test the middle of the chain: did the analyst's list have 31 people before Gmail?

Swap a known-good part

Use a sample you trust — e.g. sample-explainer.mp4 for the upload test.

Compare with the baseline

What changed since it last worked? Version history shows the exact difference.

Escalate with evidence

After three rounds, hand a person the test, the result and what you tried.

© Tertiary Infotech Academy Pte Ltd

From Skill to Schedule

Lab 13: do the job once, keep it as a skill, let a schedule run it.

01

Do

The Growth Analyst writes the weekly report once, from Drive.

02

Verify

Fix the sections until the report is right.

03

Skill

campaign-report holds the procedure.

04

Schedule

A scheduled task runs it every Monday at 8am.

05

Review

It leaves a Gmail draft — a person decides.

TIP Pause scheduled tasks you are not using.

© Tertiary Infotech Academy Pte Ltd

Automate the Preparation, Not the Decision

What a scheduled agent should and should not do for a regulated firm.

Prepare

Read the data, run the skill, draft the report, flag what needs a decision.

Draft, never send

Write "draft only — never send, publish or approve" in the skill AND the schedule.

Keep a human gate

The schedule fills an inbox. A named person still decides.

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report

LAB 13 | DAY 2 · ~30 MIN | USE: CLAUDE CODE

THE STORY SO FAR

The team must keep working after class. Rachel wants the Growth Analyst's report in her inbox every Monday — as a draft she decides on.

YOU'LL BUILD

a scheduled task that drafts the weekly growth report every Monday at 8am, and one draft in Gmail

Surface: Claude Code (skill, routine, Drive, Gmail)

Features: the Growth Analyst's campaign-report skill · do · verify · schedule · draft-only delivery

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 13". Run all of this lab's prompts in it.

2. **Do it once** Copy and paste Prompt A. Refine the report until it is right.

3. **Schedule it** After the report is right, copy and paste Prompt B into the instructions of a new routine: Routines → New routine → Local, Mondays at 8am.

4. **Run it now** Click Run now once, approve any permission prompts, then open the Gmail draft.

5. **Pause it** After class, pause the routine.

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report

PROMPT A — Claude Code: do it once

Act as Horizon's Growth Analyst. Using Google Drive, read campaign-results.csv in the Horizon Studio folder. Use the campaign-report skill and data/weekly-report-brief.md. Leave the report as a Gmail DRAFT to me, subject "Horizon weekly growth report" and today's date. Send nothing. Also publish the report as an artifact: cost per chat by channel as a chart, and the budget split. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report (cont.)

PROMPT B — Claude Code: the routine

Run the campaign-report skill as Horizon's Growth Analyst. Read campaign-results.csv in the Horizon Studio folder on Google Drive and follow weekly-report-brief.md. Leave the report as a draft email to me in Gmail with the subject "Horizon weekly growth report" and today's date. Do not send it, and do not publish or approve anything.

NOTE Automate the preparation, not the decision. The schedule fills an inbox; a person still decides.

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report (cont.)

FOLLOW-UP — type any of these in the chat

What is the one decision this report asks me to make?

Which channel gets us a chat most cheaply, and how sure are you?

What would you change in next month's budget, and why?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Campaign 5: Always On — the Weekly Growth Report

Check your work before moving on.

CHECK YOUR WORK

- ✓ The report used the campaign-report skill from your folder.
- ✓ The report names the cheapest channel per chat (the newsletter).
- ✓ It recommends a budget split with a reason per channel.
- ✓ A scheduled task runs Mondays at 8am with draft-only instructions.
- ✓ A manual run left a draft — nothing was sent.
- ✓ The task is paused after class.

© Tertiary Infotech Academy Pte Ltd

Responsible AI in the Team

Each risk has a control you built in this course.

RISK	CONTROL	WHERE
Wrong figure	fact-check skill; the Lead's review; facts sheet	Labs 3, 5, 10–12
Read as advice or a promise	fin-compliance; disclaimers; no returns	Labs 3, 10–12
An agent with too much access	Least-privilege connectors per agent	Lab 4
Unapproved post	Approval typed by a person, hash, gate	Lab 10
Marketing without consent	The Analyst filters by consent	Lab 11
Undisclosed AI media	Description line + YouTube disclosure	Lab 12

TIP Keep keys in .env, client data out of prompts, and the approval records for audit.

© Tertiary Infotech Academy Pte Ltd

Measure What Matters

Lab 14: the numbers decide next month — and which patterns earn their tokens.

Cost per chat

Spend ÷ chats booked shows which channel works: email first, YouTube last.

Defects before a person

A Lead that catches mistakes saves the reviewer's time; one agent alone caught none.

The playbook

Roster, instructions, skills, connectors, approval matrix, incidents — so the team runs without you.

© Tertiary Infotech Academy Pte Ltd

Versioning the Team

Your agents, skills and rules are an integration plan. Version it, so every change can be compared — and undone.

01

Baseline

Save the team as it works today: version 1.0, with a note.

02

Observe

Results, run log, QA rounds, incidents — what the outcomes say.

03

Propose

One change, with the evidence, the files it touches and the risk.

04

Test

Run the same checks as before: fact-check, fin-compliance, a dry run.

05

Release or roll back

Save version 1.1 — or go back to 1.0 if a check fails.

TIP Claude Code keeps version history with git for you: ask in plain English, "save this as version 1.1".

© Tertiary Infotech Academy Pte Ltd

A Good Change Proposal

Lab 14: propose one change to how the team integrates, then prove it before it becomes the new version.

What we saw

The outcome that prompts the change, with numbers from the run log or the report.

The change

One change: a rule, a skill, a connector or a pattern. Files it touches.

How we test it

The check that must still pass, and the one that should now improve.

Risk and roll back

What could break, and the version to return to.

Who approves

A named person — Rachel for anything with CPF or tax figures.

Version note

"1.1 — what changed and why", saved with the files.

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook

LAB 14 | DAY 2 · ~40 MIN | USE: CLAUDE CODE

THE STORY SO FAR

Eight weeks in. Rachel asks which ways of running agents earned their tokens, and how the team runs safely after you leave.

YOU'LL BUILD

the agent-pattern comparison, the team playbook, a change proposal, and versions 1.0 and 1.1 of the team — shown as artifacts, with a Word copy of each

Surface: Claude Code (subagents)

Features: compare agent patterns · the roster: instructions, skills, connectors · approval matrix · responsible-AI controls · version history · a change proposal and a roll-back plan

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 14". Run all of this lab's prompts in it.
2. **Check the data** The assets are already in data/. Open one and skim it.
3. **Compare the patterns** Copy and paste Prompt A.
4. **Write the playbook** After the Lead shows the comparison, copy and paste Prompt B.
5. **Review it as Rachel** Add one rule of your own and say why.
6. **Propose and version a change** After you have added your rule, copy and paste Prompt C. Claude saves version 1.0, makes one change, tests it and saves 1.1.
7. **Compare the versions** Read what changed between 1.0 and 1.1. If a check failed, ask: "Roll back to version 1.0."

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook

PROMPT A — Claude Code: compare

Use the growth-analyst subagent. From data/run-log.csv compare the patterns we used — one agent, subagents, agent teams and the scheduled task — on minutes, tokens, defects caught before a person and human edits. Which pattern pays off for which job? Save reports/agent-comparison.md. Publish it as an artifact: one chart per measure, comparing the patterns. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook (cont.)

PROMPT B — Claude Code: the playbook

Write strategy/team-playbook.md with data/playbook-outline.md. Include the roster: for each agent its instructions file in .claude/agents/, its skills, connectors, model and what it may never do; the approval matrix by risk; the rules in data/responsible-ai-checklist.md; and what to do when a wrong post goes out. Then run fin-compliance on it. Publish the playbook as an artifact: the roster as cards, the approval matrix as a table. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook (cont.)

PROMPT C — Claude Code: version and propose a change

Keep a version history of this folder (use git; set it up if it is not there yet). Save everything as version 1.0 with the note "The team as built in class". From reports/agent-comparison.md and data/run-log.csv, propose ONE change to how the team works. Write review/change-proposal.md: what we saw, the change, the files it touches, how we test it, the risk, who approves, and how we roll back. Make the change, test it on one recent piece with fact-check and fin-compliance, and save version 1.1 with a one-line note. Show me what changed between
1.0 and 1.1, and how to go back to 1.0. Publish the proposal and the version history as an artifact. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Teams cost more tokens. The run log shows what you buy with them: defects caught before a person, and fewer edits.

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook (cont.)

FOLLOW-UP — type any of these in the chat

In one line each: when should we use one agent, subagents and an agent team?

Who must approve a post that has a CPF figure in it?

A wrong post went out. What are the first three things we do?

After Prompt C: In plain English, what changed in version 1.1, and how would we undo it?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Measure, Version and Write the Team Playbook

Check your work before moving on.

CHECK YOUR WORK

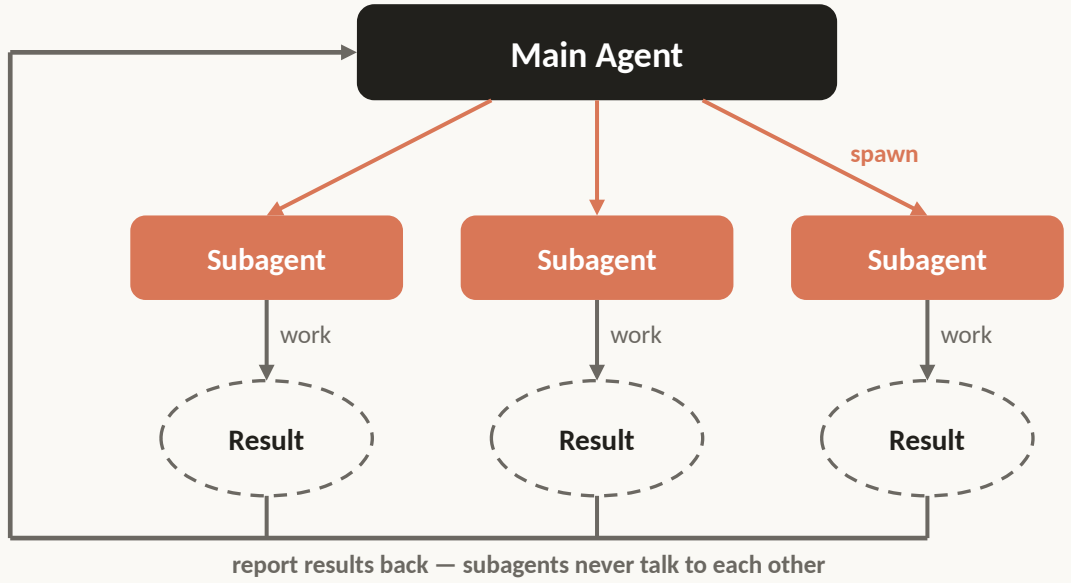
- ✓ The comparison shows the one-agent baseline caught 0 defects.
- ✓ It says which pattern suits which job, with numbers.
- ✓ The roster lists all six agents with skills and connectors.
- ✓ The approval matrix says who approves CPF and tax figures.
- ✓ It has an incident procedure, and your own rule.
- ✓ Versions 1.0 and 1.1 are saved, each with a note.
- ✓ The change proposal has the evidence, the test, the risk and the roll-back.

© Tertiary Infotech Academy Pte Ltd

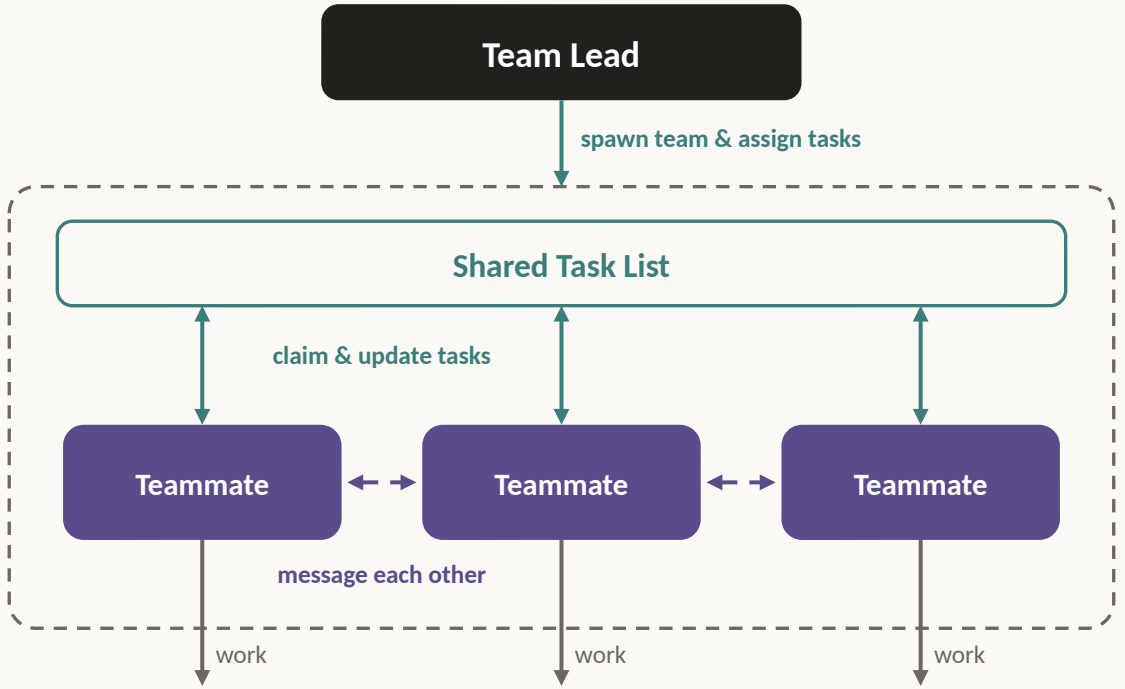
Subagents vs Agent Teams

Same work, different wiring: who spawns whom, and who talks to whom.

SUBAGENTS



AGENT TEAM



Subagents report back to one agent. Teammates share a task list and talk to each other.

© Tertiary Infotech Academy Pte Ltd

A Team That Argues

Lab 15, Part 1: test a strategy before Horizon spends on it.

Advocate · growth-strategist

- Argues FOR the "Turning 55" CPF explainer
- Cites results, personas and the facts sheet
- Answers the critic directly — two rebuttals

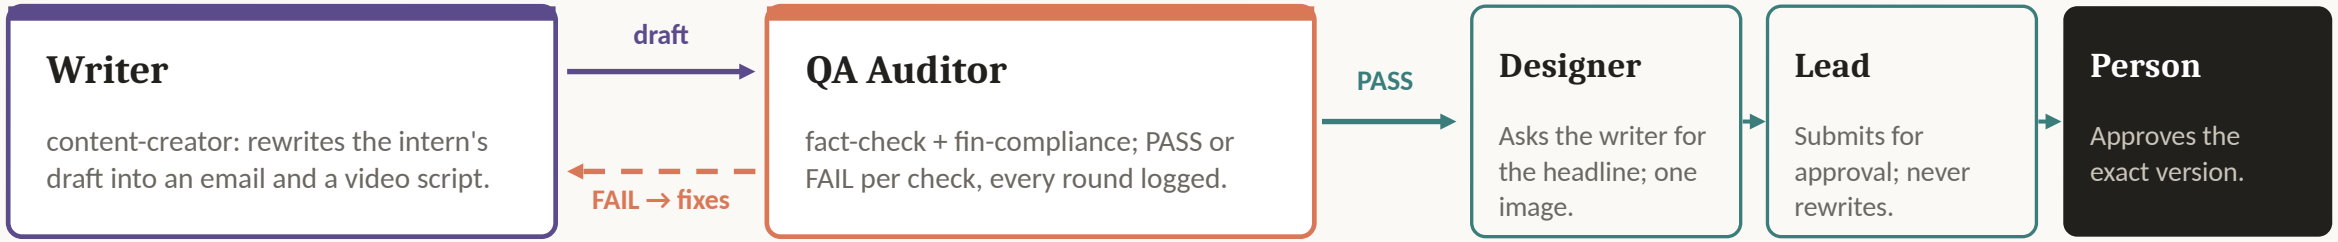
Critic · growth-analyst

- Argues AGAINST it — cost, timing, audience
- Every point backed by a file, not an opinion
- The Lead judges; a person makes the call

© Tertiary Infotech Academy Pte Ltd

The QA Loop

Lab 15, Part 2: a failed audit goes straight back to the writer — not through the Lead.



THE AUDIT FAILS WHEN

1. A figure is not in facts-2026.md
2. A promise of savings or returns
3. Advice ("everyone should...")
4. A wrong age or a missing disclaimer

Three rounds at most — then the Lead escalates to a person.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop

LAB 15 | DAY 2 · ~25 MIN | USE: CLAUDE CODE

THE STORY SO FAR

January is for people turning 55. Before Horizon spends on a CPF explainer for them, Rachel wants it argued both ways — and the intern's draft checked until it passes.

YOU'LL BUILD

the decision (pros, cons, verdict), then an email, a video script and a thumbnail that passed QA, submitted for your approval — shown as artifacts, with a Word copy of each

Surface: Claude Code (agent teams: Lead, teammates, shared task list, direct messages)

Features: an advocate and a critic argue the pros and cons · the Lead judges · a writer and a QA Auditor loop until PASS · FAIL goes straight back to the writer · human approval

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

1. **New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 15". Run all of this lab's prompts in it.
2. **Part 1: the debate** Copy and paste Prompt A. The advocate and the critic message each other; the decision page shows their debate.
3. **Decide** Read the decision artifact. Reply "go", "go with conditions" or "no-go" — the call is yours.
4. **Part 2: the QA loop** After you have made the call, copy and paste Prompt B. The audit fails on round 1; watch the auditor message the writer directly.
5. **Approve as a person** Read the pieces and the QA artifact, then type in the chat: approved <id> by <your name> (or: changes <id>: what to fix).
6. **Shut down** After you have approved the pieces, copy and paste Prompt C. The Lead stops each teammate and cleans up the team.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop

PROMPT A — Claude Code: the strategy debate

Create an agent team to test the strategy in data/turning-55-brief.md before we spend on it. You are the Lead and the judge: do not argue. Spawn two teammates:
- advocate, using the agent type growth-strategist:
  argue FOR the strategy.
- critic, using the agent type growth-analyst:
  argue AGAINST it.
- Both: cite data/campaign-results.csv, strategy/personas.md, data/facts-2026.md or data/compliance-checklist.md for every point; message each other directly; rebut twice at most.
- You: write strategy/turning-55-decision.md with the pros, the cons, your verdict (go / go with conditions / no-go) and the conditions. Publish the decision as an artifact with the verdict on top, the pros and cons in two columns, the debate (the messages the advocate and critic sent each other, in order) and what I do next. Then wait for my decision. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop (cont.)

PROMPT B — Claude Code: the QA loop

Clean up the debate team, then create a new agent team for Part 2 of data/turning-55-brief.md, using my decision. Spawn three teammates:
- writer, agent type content-creator: rewrite data/turning-55-draft.md into content/turning-55/email.md and content/turning-55/video-script.md.
- qa-auditor: audit each piece with fact-check and fin-compliance; write review/turning-55-qa.md, one section per round, PASS or FAIL per check. On FAIL, message the writer directly with the lines to fix and audit again. Three rounds at most.
- designer, agent type creative-designer: once a piece passes, ask the writer for its headline and make one thumbnail with alt text.
- You, the Lead: never write or fix content. When everything passes, submit each piece with scripts/submit.mjs. Publish nothing. Publish the QA report as an artifact, one row per round, PASS or FAIL per check, in colour. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop (cont.)

PROMPT C — Claude Code: shut down

Ask every teammate to shut down, then clean up the team. Show me the final task list and how many QA rounds each piece needed. Publish it as an artifact.

NOTE Use subagents when only the result matters; use a team when the agents must talk — to argue, or to send work back. Every teammate is its own Claude session, so teams cost more tokens.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop (cont.)

FOLLOW-UP — type any of these in the chat

Summarise the debate: the best point for, the best point against, and your verdict.

What failed in the first QA round, in plain English?

How many rounds did each piece need, and why?

NOTE Ask until the result makes sense to you — then decide.

© Tertiary Infotech Academy Pte Ltd

Agent Teams: A Strategy Debate and a QA Loop

Check your work before moving on.

CHECK YOUR WORK

- ✓ The advocate and the critic messaged each other; every point cites a file.
- ✓ The decision has pros, cons, a verdict — and you made the call.
- ✓ Round 1 of the QA report FAILED: the FRS figure, CPF LIFE "from 55", the promised payout, the advice and the missing disclaimer.
- ✓ The auditor sent the fixes to the writer directly; a later round PASSED.
- ✓ You approved by typing in the chat; the rows show your name and a hash.
- ✓ Both teams were shut down and cleaned up.

© Tertiary Infotech Academy Pte Ltd

Integration Best Practices

Ten habits from the course — each one a control you built.

PRACTICE	WHY	LAB
Assess compatibility first	A connection can work and still be wrong	Lab 4
Least privilege per agent	An agent cannot misuse what it cannot reach	Lab 4
Draft, never send	A person sends and publishes	Labs 9–13
Dry run before live	See the exact request first	Labs 10, 12
A human gate a hook enforces	A prompt persuades; a hook holds	Lab 10
Test with known answers	31 of 40, under 90 s, facts-sheet figures	Labs 6–12
Version, change one thing, roll back	Every change can be compared and undone	Lab 14

TIP Also: keys only in .env, client data out of prompts, pause routines you are not using.

© Tertiary Infotech Academy Pte Ltd

How Each Pattern Is Triggered

Five ways your team runs. The difference is what sets each one off.

BY DELEGATION

Subagents

The Lead starts them when you ask — or when a task matches an agent's description.

e.g.

```
"use subagents,
one per role"
```

.claude/agents/

BY THE LEAD

Agent teams

The Lead spawns teammates who share a task list and message each other.

e.g.

```
"spawn teammates
using agent types…"
```

Claude Code (experimental)

ON DEMAND

Skills

Loaded when the task matches the description. Free until then.

e.g.

```
campaign-brief
video-render
```

.claude/skills/

ON AN EVENT

Hooks

Run every time an event fires — before a tool runs. The model cannot skip them.

e.g.

```
PreToolUse:
gate-hook.mjs
```

.claude/settings.json

ON THE CLOCK

Scheduled tasks

Run a skill on a schedule and leave the result for a person.

e.g.

```
Mondays 8am:
campaign-report
```

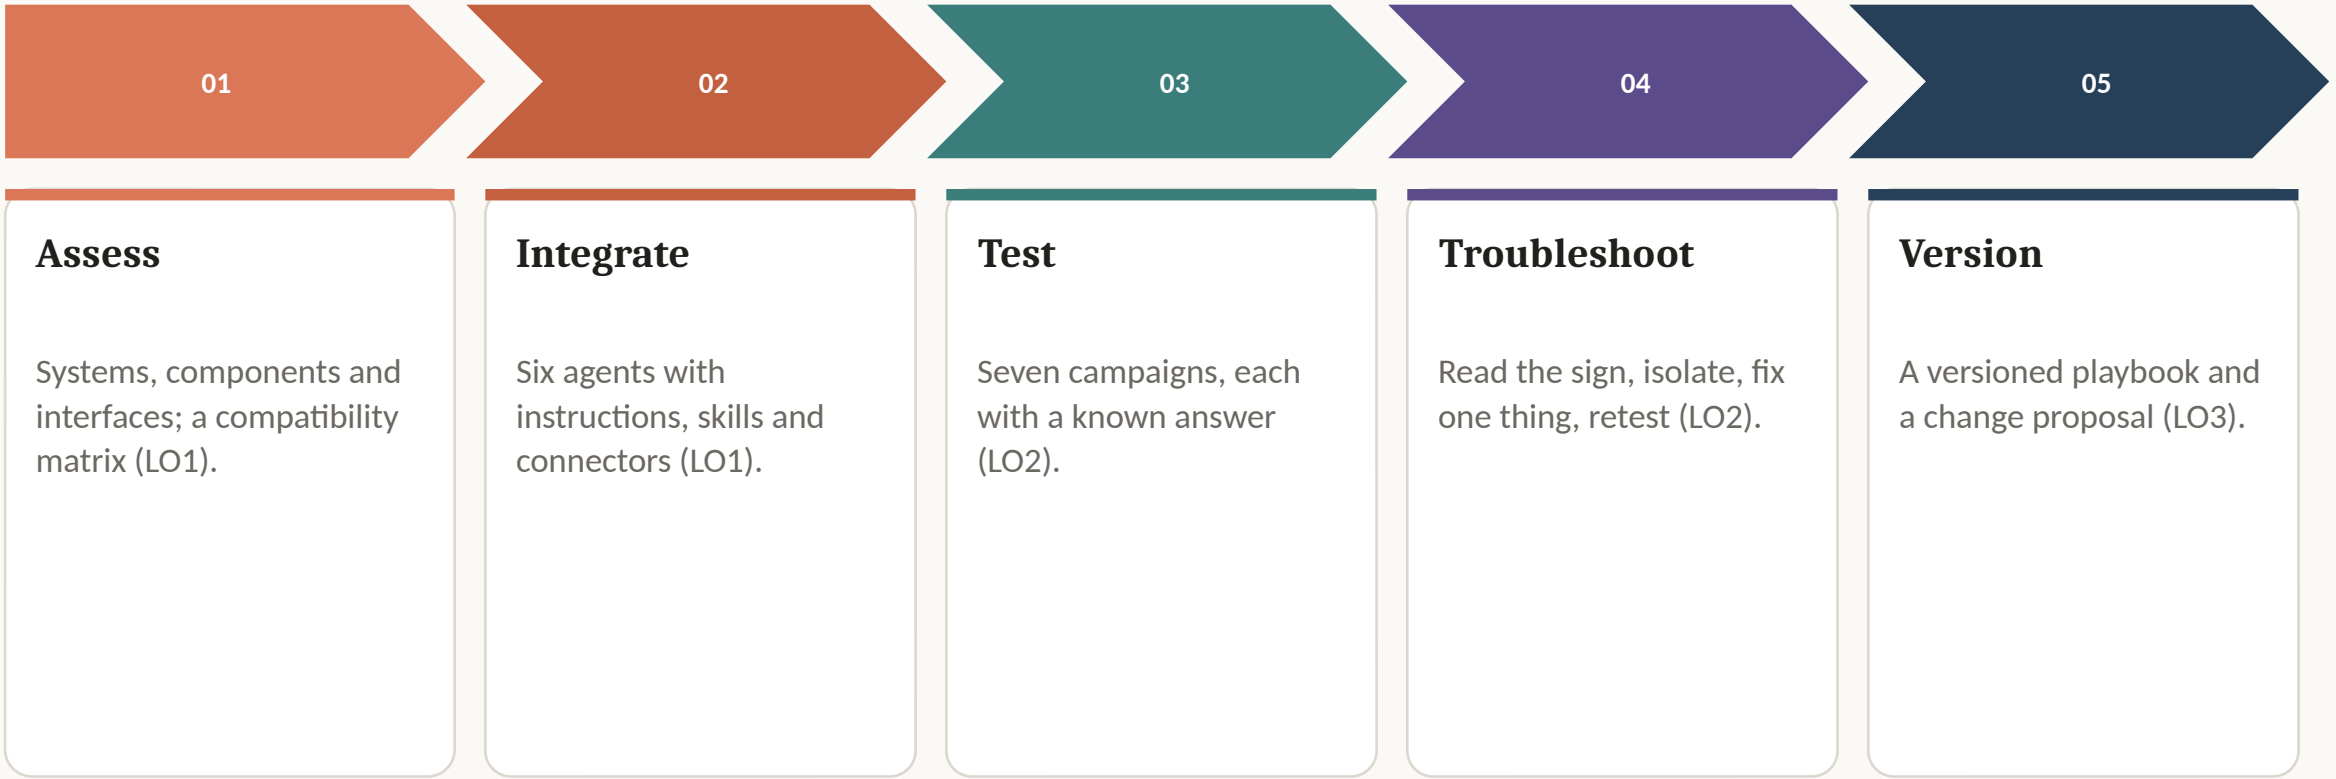
Claude

Who decides? You: the approval · the model: delegation and skills · the system, every time: hooks and the clock

© Tertiary Infotech Academy Pte Ltd

Horizon: One Team, Integrated and Tested

Two days, one regulated business, three learning outcomes.



© Tertiary Infotech Academy Pte Ltd

LEARNING MANAGEMENT SYSTEM

Courseware & Assessment on the LMS

Log in

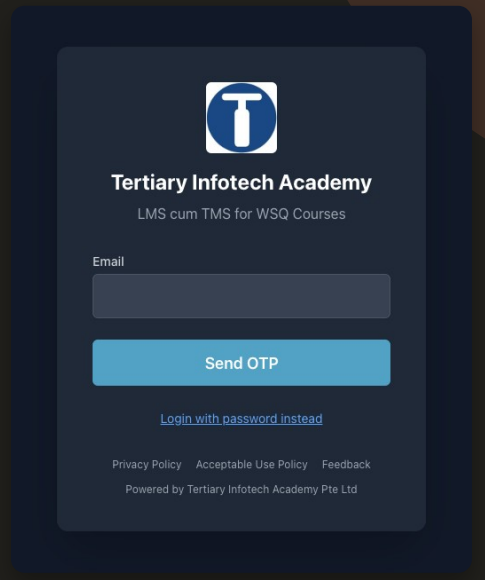
Use your enrolled account to access the course materials and assessment.

Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.

Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



LMS: https://lms-tms.tertiaryinfotech.com/

© Tertiary Infotech Academy Pte Ltd

BEFORE YOU GO

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

https://lms-tms.tertiaryinfotech.com/



Digital Attendance

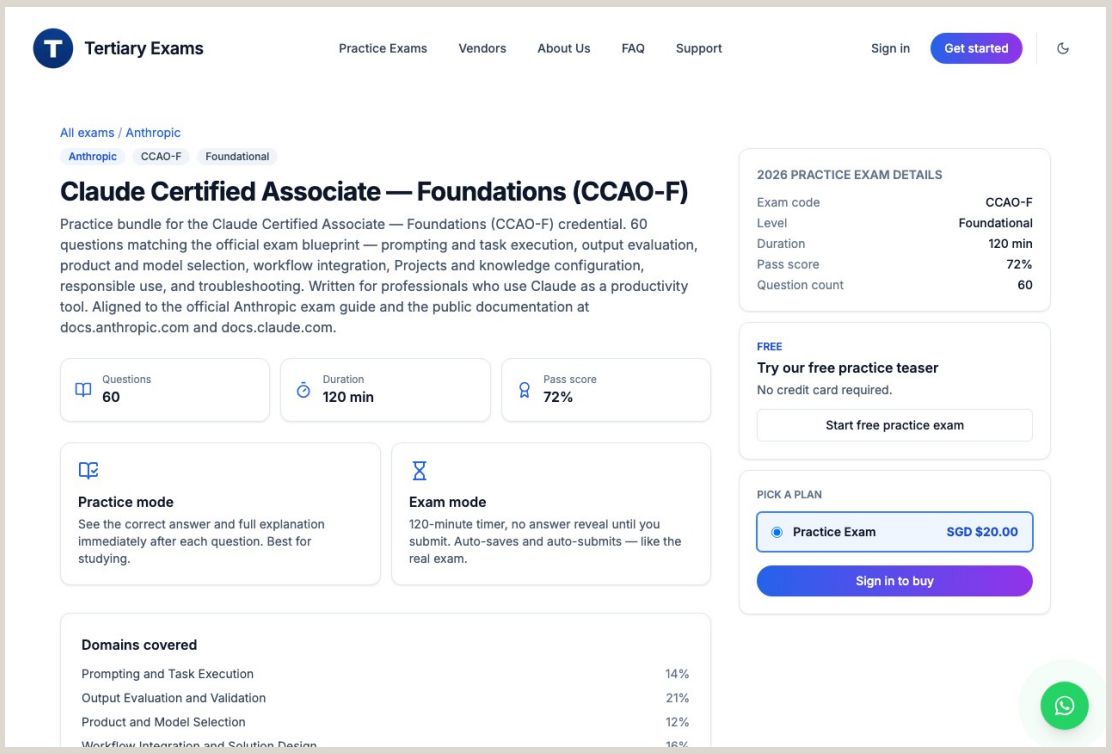
Remember to take AM, PM and Assessment digital attendance — required for your funding.

© Tertiary Infotech Academy Pte Ltd

AFTER THE COURSE

Practice Exam: Claude Certified Associate

exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations



PRACTICE EXAM exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations

CCAO-F · Foundations

For professionals who use Claude as a productivity tool.

What it covers

Prompting and task execution, output evaluation, model selection, workflow integration, responsible use.

Try it free

Start with the free practice teaser; exam mode is 60 questions in 120 minutes.

© Tertiary Infotech Academy Pte Ltd

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 3:30 PM
- Practical Performance (PP) — 1.5 hours
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance

Based on SSG digital attendance records.

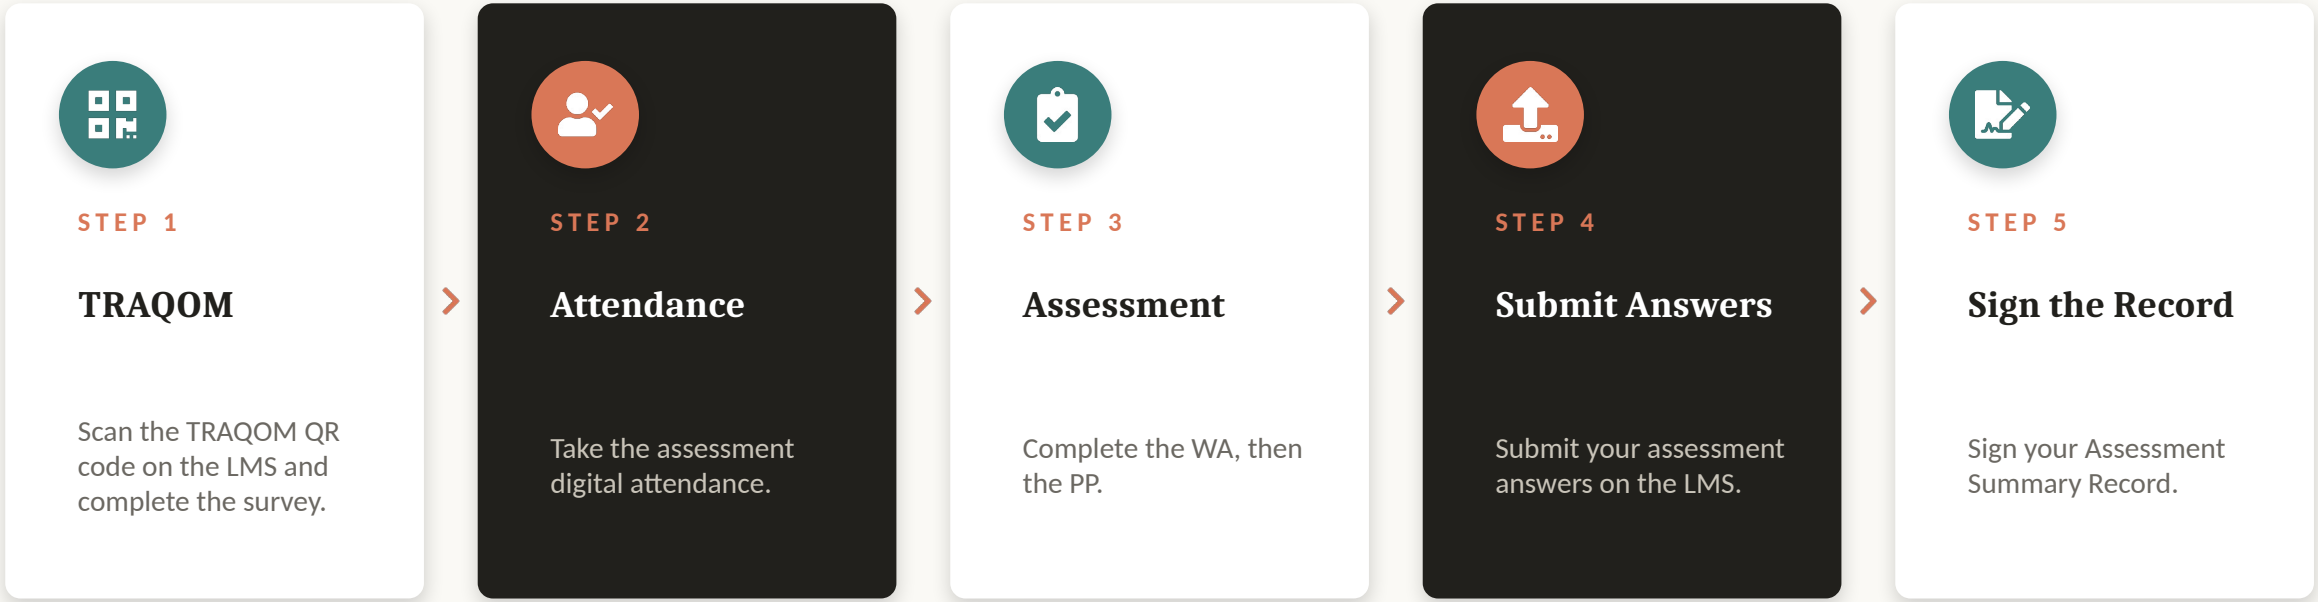
Assessed as ‘Competent’

Complete and pass both assessment components.

© Tertiary Infotech Academy Pte Ltd

Assessment Flow

At the assessment stage, follow these five steps in order.



TIP *Courseware and the assessment are on the LMS: https://lms-tms.tertiaryinfotech.com/*

© Tertiary Infotech Academy Pte Ltd

HOUSEKEEPING

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

Minimum 75% attendance required

© Tertiary Infotech Academy Pte Ltd

Thank You!

Questions? We're here to help — during and after class.

enquiry@tertiaryinfotech.com | +65 6100 0613